



3D Point Clouds

Lecture 5 – Deep Learning with Point Clouds

主讲人 黎嘉信

Aptiv 自动驾驶
新加坡国立大学 博士
清华大学 本科





1. Introduction to Deep Learning



2. PointNet



3. PointNet++



Intended Learning Outcomes

1

Explain the principles of the operations of different layers and training algorithms

2

Compare different neural network architectures

3

Implement popular neural networks

4

Solve problems using neural networks and deep learning techniques

ARTIFICIAL INTELLIGENCE

IS NOT NEW

ARTIFICIAL INTELLIGENCE

Any technique which enables computers to mimic human behavior



1950's

1960's

1970's

1980's

MACHINE LEARNING

AI techniques that give computers the ability to learn without being explicitly programmed to do so



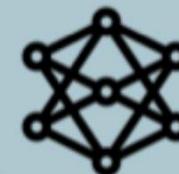
1990's

2000's

2010s

DEEP LEARNING

A subset of ML which make the computation of multi-layer neural networks feasible



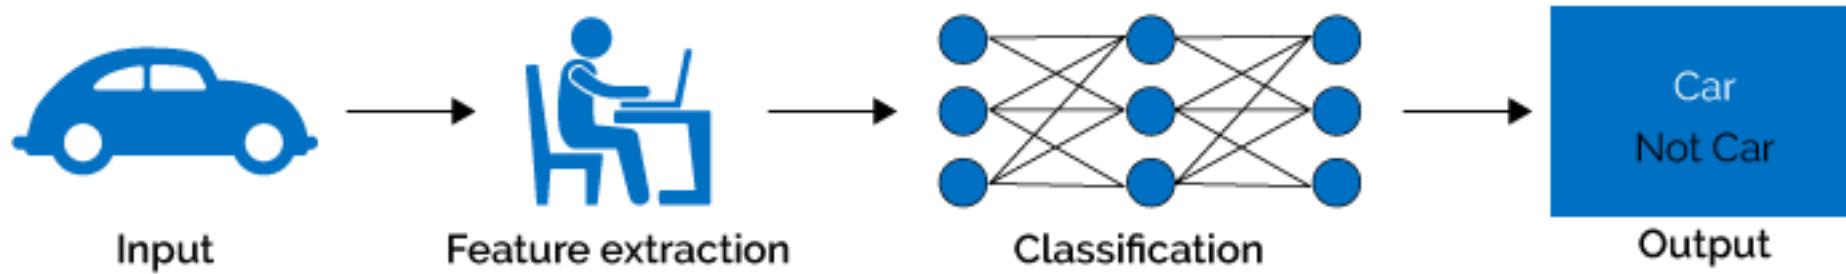
ORACLE

Copyright © 2018, Oracle and/or its affiliates. All rights reserved. |

<https://blogs.oracle.com/bigdata/difference-ai-machine-learning-deep-learning>



Machine Learning



Deep Learning



Feature extraction + classification



Everything is an optimization problem.

– Stephen Boyd

We can't really solve most optimization problems.

Deep Learning saves the world because it is a
magical optimization solver.

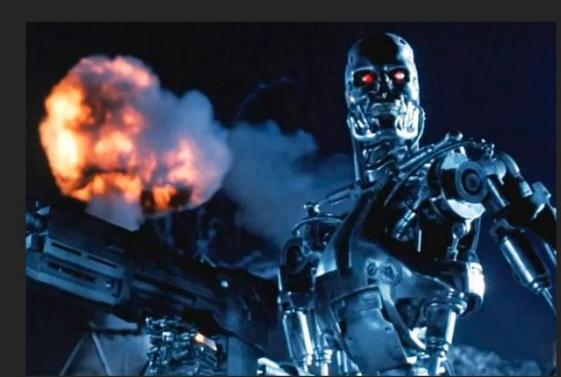
Stephen Boyd and
Lieven Vandenberghe

Convex Optimization

CAMBRIDGE



Another explain of Deep Learning



What society thinks I do



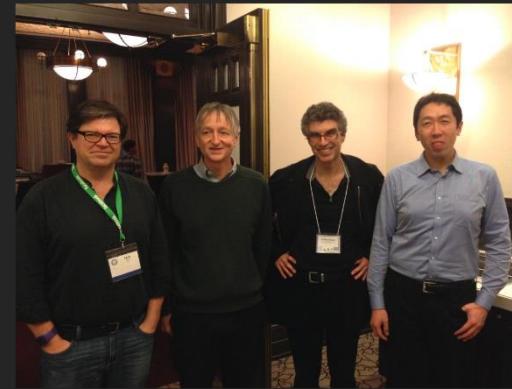
What my friends think I do



What other computer
scientists think I do



What mathematicians think I do



What I think I do

```
from keras import *
```

What I actually do



- Deep learning is trying to solve one problem

$$\min_x f(x)$$

- Let's start from linear function

$$y = f(x) = w^T x + b$$

- Given a set of $\{x_i \in \mathbb{R}^n, y_i \in \mathbb{R}\}$
- Solve $w \in \mathbb{R}^n, b \in \mathbb{R}$



Consider the problem of house price prediction

A dataset consist of M instance

- Features $x_i \in \mathbb{R}^n$
- Ground truth (price) $y_i \in \mathbb{R}$

Unknown parameter $w \in \mathbb{R}^n, b \in \mathbb{R}$

Denote prediction of one instance

$$\hat{y}_i = w^T x_i + b = \begin{bmatrix} w \\ b \end{bmatrix} [x_i \quad 1]$$

A bit abuse of notation with *homogeneous coordinate*

$$\hat{y}_i = w^T x_i$$

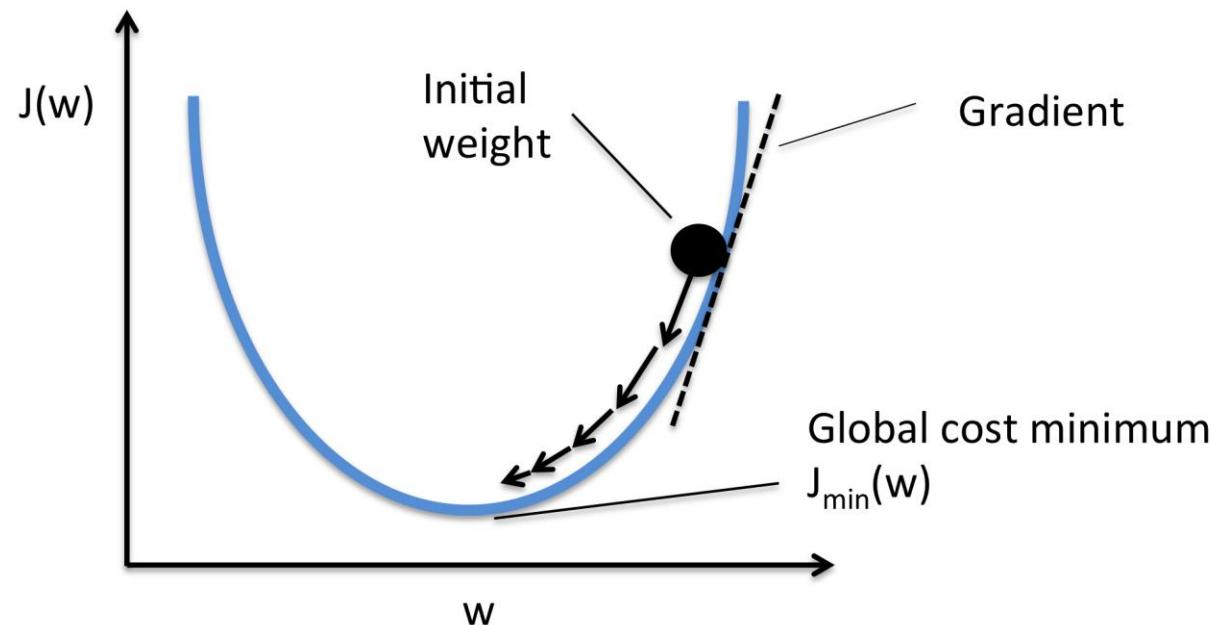
$$x_i = \begin{bmatrix} x_{i1} \\ x_{i2} \\ x_{i3} \\ \vdots \\ x_{in} \end{bmatrix} = \begin{bmatrix} size \\ floor \\ age \\ \vdots \\ * * * \end{bmatrix}$$



- It is a linear regression problem

$$w = \arg \min_w \sum_{i=1}^M \frac{1}{2} (w^T x_i - y_i)^2$$

- Typical $Ax = b$ problem, but can we solve it by *Gradient Descent*?





Gradient Descent

- For any differentiable function $F(x): \mathbb{R}^n \rightarrow \mathbb{R}^m$
- The objective is to minimize $F(x)$

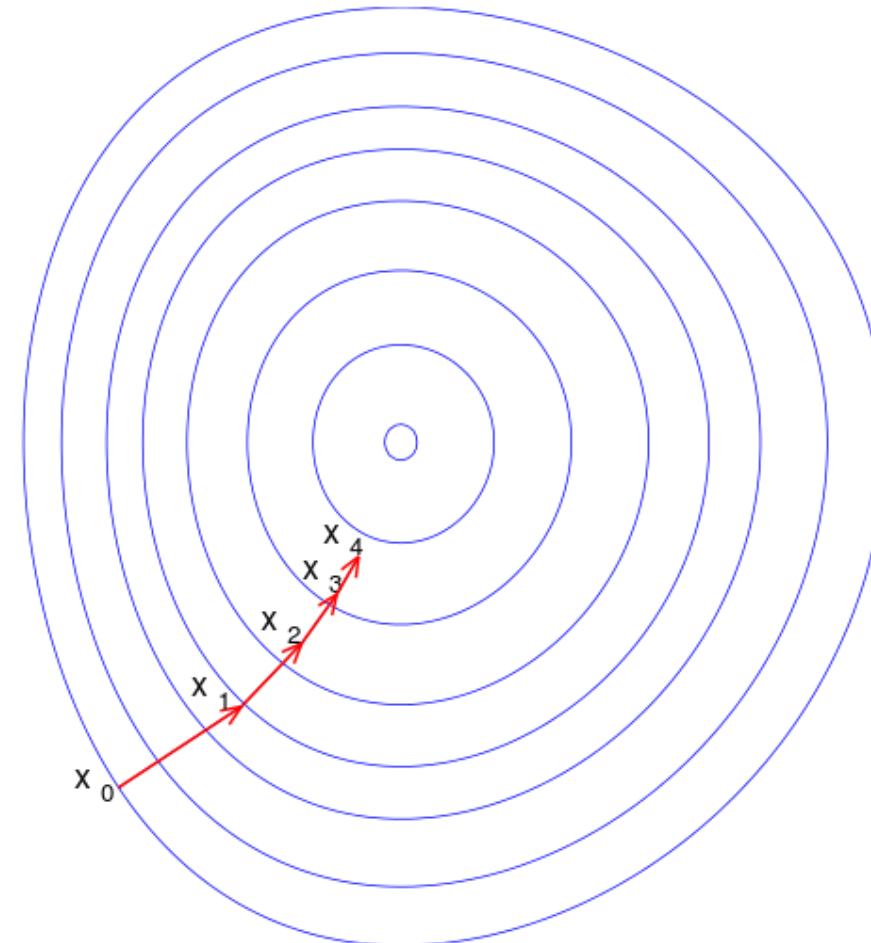
$$x^* = \operatorname{argmin}_x F(x)$$

- At position x_n , $F(x_{n+1}) \leq F(x_n)$ is guaranteed if,

$$x_{n+1} = x_n - \gamma \nabla F(x_n)$$

where $\nabla F(x_n)$ is the gradient of F at x_n

- γ is the *step size*





Objectives

$$w = \arg \min_w \sum_{i=1}^M \frac{1}{2} (w^T x_i - y_i)^2$$

Consider a single instance

$$J(w) = \frac{1}{2} (w^T x - y)^2 = \frac{1}{2} z^2, \quad z = w^T x - y$$

Gradient descent requires

$$\frac{\partial J(w)}{\partial w} = \frac{\partial J}{\partial z} \frac{\partial z}{\partial w} = z x = (w^T x - y)x$$

where $J(w) \in \mathbb{R}$, $z \in \mathbb{R}$, $w \in \mathbb{R}^n$



Constant Rule $f(x) = c \rightarrow f'(x) = 0$

Constant Multiple Rule $f(x) = c \cdot g(x) \rightarrow f'(x) = cg'(x)$

Power Rule $f(x) = x^n \rightarrow f'(x) = nx^{n-1}$

Sum Rule $f(x) = g(x) + h(x) \rightarrow f'(x) = g'(x) + h'(x)$

Product Rule $f(x) = g(x)h(x) \rightarrow f'(x) = g'(x)h(x) + g(x)h'(x)$

Quotient Rule $f(x) = \frac{g(x)}{h(x)} \rightarrow f'(x) = \frac{g'(x)h(x) - g(x)h'(x)}{h(x)^2}$

Chain Rule $f(x) = g(u), u = h(x) \rightarrow f'(x) = g'(u)h'(x)$



- Given a multivariable function $f(x, y)$, and two univariate functions $x(t), y(t)$, the **multivariable chain rule** is:

$$\frac{df(\textcolor{blue}{x}(t), \textcolor{red}{y}(t))}{dt} = \frac{\partial f}{\partial \textcolor{blue}{x}} \frac{dx}{dt} + \frac{\partial f}{\partial \textcolor{red}{y}} \frac{dy}{dt}$$

- How to compute $\frac{\partial f(x)}{\partial x}$ if $f(x)$ and/or x is scalar, vector, matrix?



$$\mathbf{x} \in \mathbb{R}^n, \mathbf{y} \in \mathbb{R}^m, \mathbf{X} \in \mathbb{R}^{n \times m}, \mathbf{Y} \in \mathbb{R}^{m \times n}$$

what are $\frac{\partial \mathbf{y}}{\partial \mathbf{x}}, \frac{\partial \mathbf{y}}{\partial \mathbf{x}}, \frac{\partial \mathbf{y}}{\partial \mathbf{x}}, \frac{\partial \mathbf{y}}{\partial \mathbf{X}}, \frac{\partial \mathbf{Y}}{\partial \mathbf{x}}$

Types	Scalar	Vector	Matrix
Scalar	$\frac{\partial y}{\partial x}$	$\frac{\partial \mathbf{y}}{\partial x}$	$\frac{\partial \mathbf{Y}}{\partial x}$
Vector	$\frac{\partial y}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$	
Matrix	$\frac{\partial y}{\partial \mathbf{X}}$		

✓ Denominator layout

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \in \mathbb{R}^{1 \times m}, \frac{\partial y}{\partial \mathbf{x}} = \mathbb{R}^n$$

Numerator layout

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} \in \mathbb{R}^m, \frac{\partial y}{\partial \mathbf{x}} = \mathbb{R}^{1 \times n}$$



$$\mathbf{x} \in \mathbb{R}^n, \mathbf{y} \in \mathbb{R}^m, \mathbf{X} \in \mathbb{R}^{n \times m}, \mathbf{Y} \in \mathbb{R}^{m \times n}$$

$$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y}{\partial x_1} \\ \frac{\partial y}{\partial x_2} \\ \vdots \\ \frac{\partial y}{\partial x_n} \end{bmatrix} \quad \frac{\partial \mathbf{y}}{\partial x} = \begin{bmatrix} \frac{\partial y_1}{\partial x} & \frac{\partial y_2}{\partial x} & \cdots & \frac{\partial y_m}{\partial x} \end{bmatrix} \quad \frac{\partial \mathbf{y}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \frac{\partial y_2}{\partial x_1} & \cdots & \frac{\partial y_m}{\partial x_1} \\ \frac{\partial y_1}{\partial x_2} & \frac{\partial y_2}{\partial x_2} & \cdots & \frac{\partial y_m}{\partial x_2} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_1}{\partial x_n} & \frac{\partial y_2}{\partial x_n} & \cdots & \frac{\partial y_m}{\partial x_n} \end{bmatrix}$$

$$\frac{\partial \mathbf{y}}{\partial \mathbf{X}} = \begin{bmatrix} \frac{\partial y}{\partial x_{11}} & \frac{\partial y}{\partial x_{12}} & \cdots & \frac{\partial y}{\partial x_{1m}} \\ \frac{\partial y}{\partial x_{21}} & \frac{\partial y}{\partial x_{22}} & \cdots & \frac{\partial y}{\partial x_{2m}} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y}{\partial x_{n1}} & \frac{\partial y}{\partial x_{n2}} & \cdots & \frac{\partial y}{\partial x_{nm}} \end{bmatrix} \quad \frac{\partial \mathbf{Y}}{\partial x} = \begin{bmatrix} \frac{\partial y_{11}}{\partial x} & \frac{\partial y_{21}}{\partial x} & \cdots & \frac{\partial y_{m1}}{\partial x} \\ \frac{\partial y_{12}}{\partial x} & \frac{\partial y_{22}}{\partial x} & \cdots & \frac{\partial y_{m2}}{\partial x} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial y_{1n}}{\partial x} & \frac{\partial y_{2n}}{\partial x} & \cdots & \frac{\partial y_{mn}}{\partial x} \end{bmatrix}$$



$\mathbf{y} =$	\mathbf{a}	\mathbf{x}	$A\mathbf{x}$	$\mathbf{x}^T A$	
$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} =$	$\mathbf{0}$	I	A^T	A	
$\mathbf{y} =$	$a\mathbf{u}$ $\mathbf{u} = \mathbf{u}(\mathbf{x})$	$v\mathbf{u}$ $v = v(\mathbf{x}), \mathbf{u} = \mathbf{u}(\mathbf{x})$	$\mathbf{v} + \mathbf{u}$ $\mathbf{v} = v(\mathbf{x}), \mathbf{u} = \mathbf{u}(\mathbf{x})$	$A\mathbf{u}$ $\mathbf{u} = \mathbf{u}(\mathbf{x})$	$g(\mathbf{u})$ $\mathbf{u} = \mathbf{u}(\mathbf{x})$
$\frac{\partial \mathbf{y}}{\partial \mathbf{x}} =$	$a \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	$v \frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \frac{\partial v}{\partial \mathbf{x}} \mathbf{u}^T$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \frac{\partial \mathbf{v}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} A^T$	$\frac{\partial g(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$
$y =$	a	$\mathbf{u}^T \mathbf{v}$ $\mathbf{v} = \mathbf{v}(\mathbf{x}), \mathbf{u} = \mathbf{u}(\mathbf{x})$	$g(u)$ $u = u(\mathbf{x})$		
$\frac{\partial y}{\partial \mathbf{x}} =$	$\mathbf{0}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \mathbf{v} + \frac{\partial \mathbf{v}}{\partial \mathbf{x}} \mathbf{u}$	$\frac{\partial g(u)}{\partial u} \frac{\partial u}{\partial \mathbf{x}}$		



Objectives

$$w = \arg \min_w \sum_{i=1}^M \frac{1}{2} (w^T x_i - y_i)^2$$

Consider a single instance

$$J(w) = \frac{1}{2} (w^T x - y)^2 = \frac{1}{2} z^2, \quad z = w^T x - y$$

Gradient descent requires

$$\frac{\partial J(w)}{\partial w} = \frac{\partial J}{\partial z} \frac{\partial z}{\partial w} = \cancel{z} \cancel{x} = (w^T x - y)x$$

where $J(w) \in \mathbb{R}$, $z \in \mathbb{R}$, $w \in \mathbb{R}^n$

Update w

$$w_{n+1} = w_n - \lambda \frac{\partial J(w)}{\partial w} \Big|_{w=w_n} = w_n - \lambda (w_n^T x - y)x$$



• What we solve just now, is

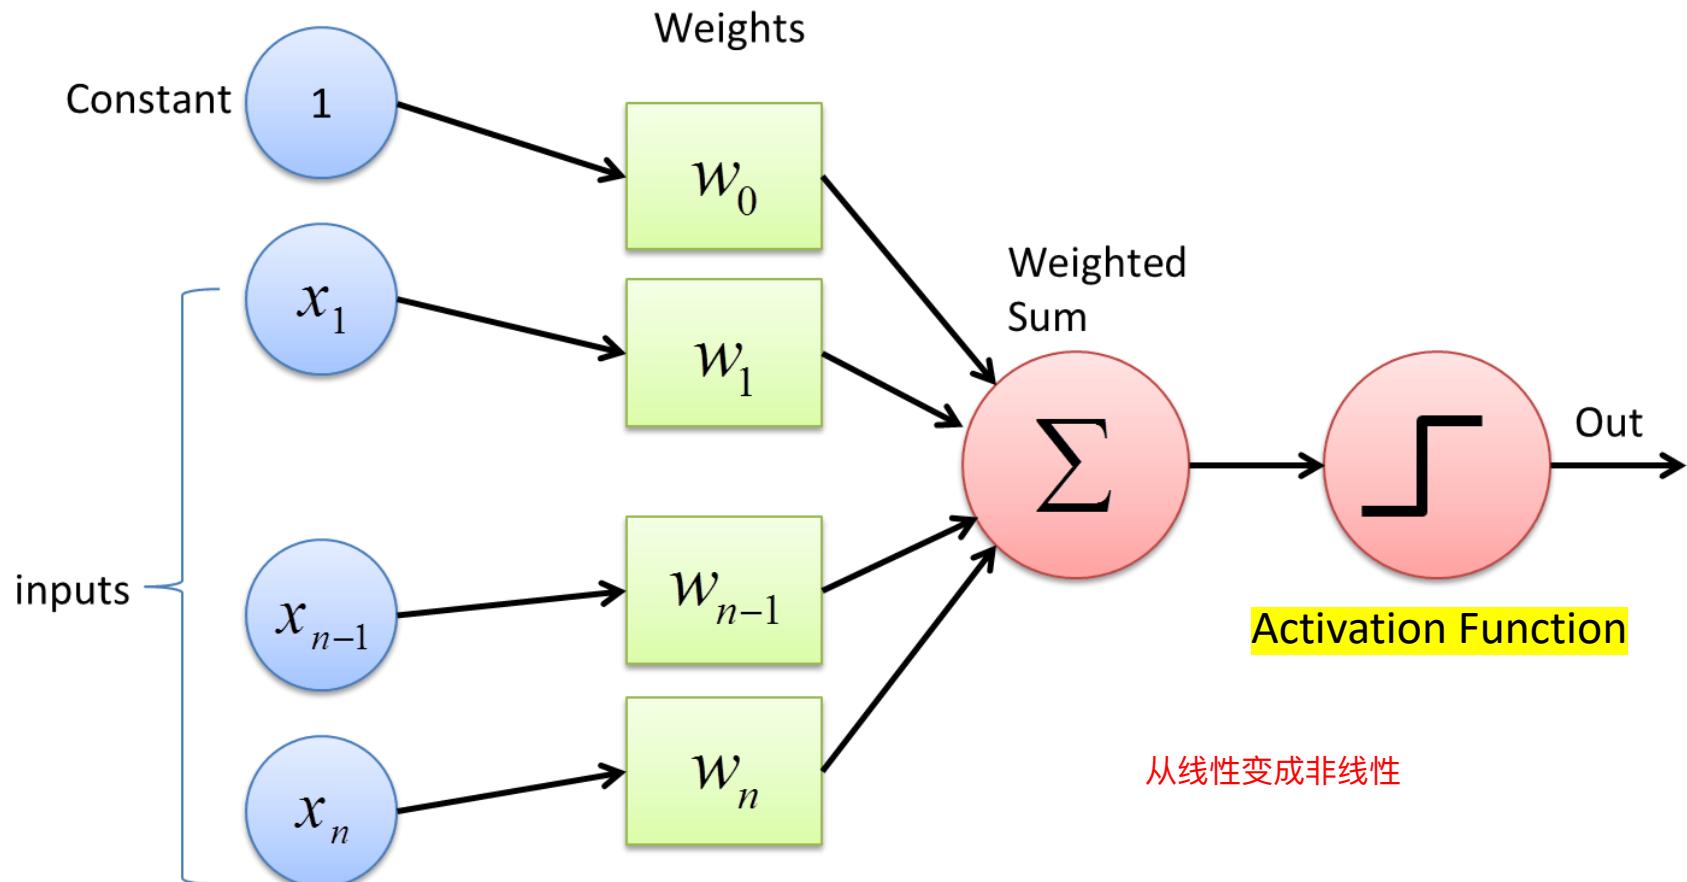
$$w = \arg \min_w \sum_{i=1}^M \frac{1}{2} (w^T x_i - y_i)^2$$

• The is the very basic Neural Network – Perceptron

- w is the trainable parameter
- x_i is the data point, y_i is the label
- Invented in 1968
- Frank Rosenblatt, Cornell Aeronautical Laboratory



Perceptron



The most simple Neural Network (NN): A Perceptron: $y = w^T x$



Train a Perceptron

Given several examples

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$$

and a perceptron

$$y = w^T x$$

Modify weight w such that \hat{y} gets ‘closer’ to y

↑
perceptron
parameter

↑
perceptron
output

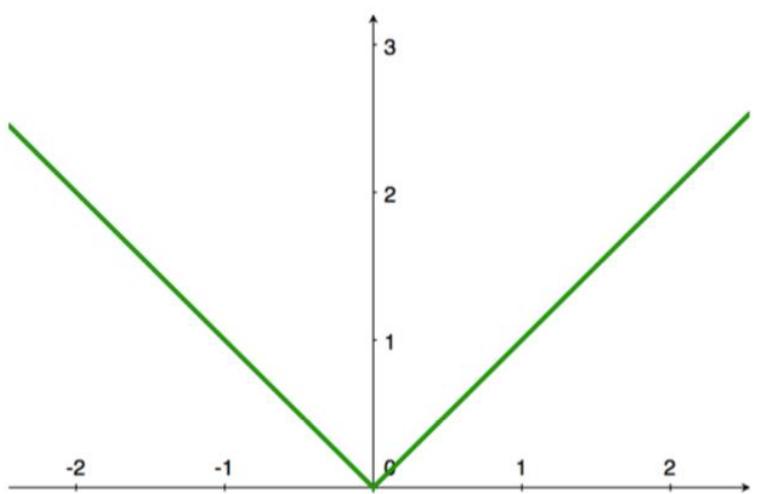
what does
this mean?

↑
true
label



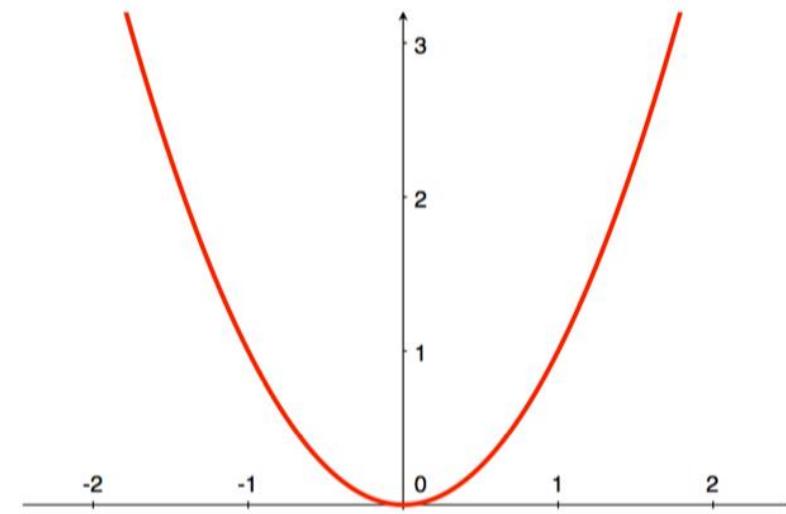
L1 Loss

$$\ell(\hat{y}, y) = |\hat{y} - y|$$



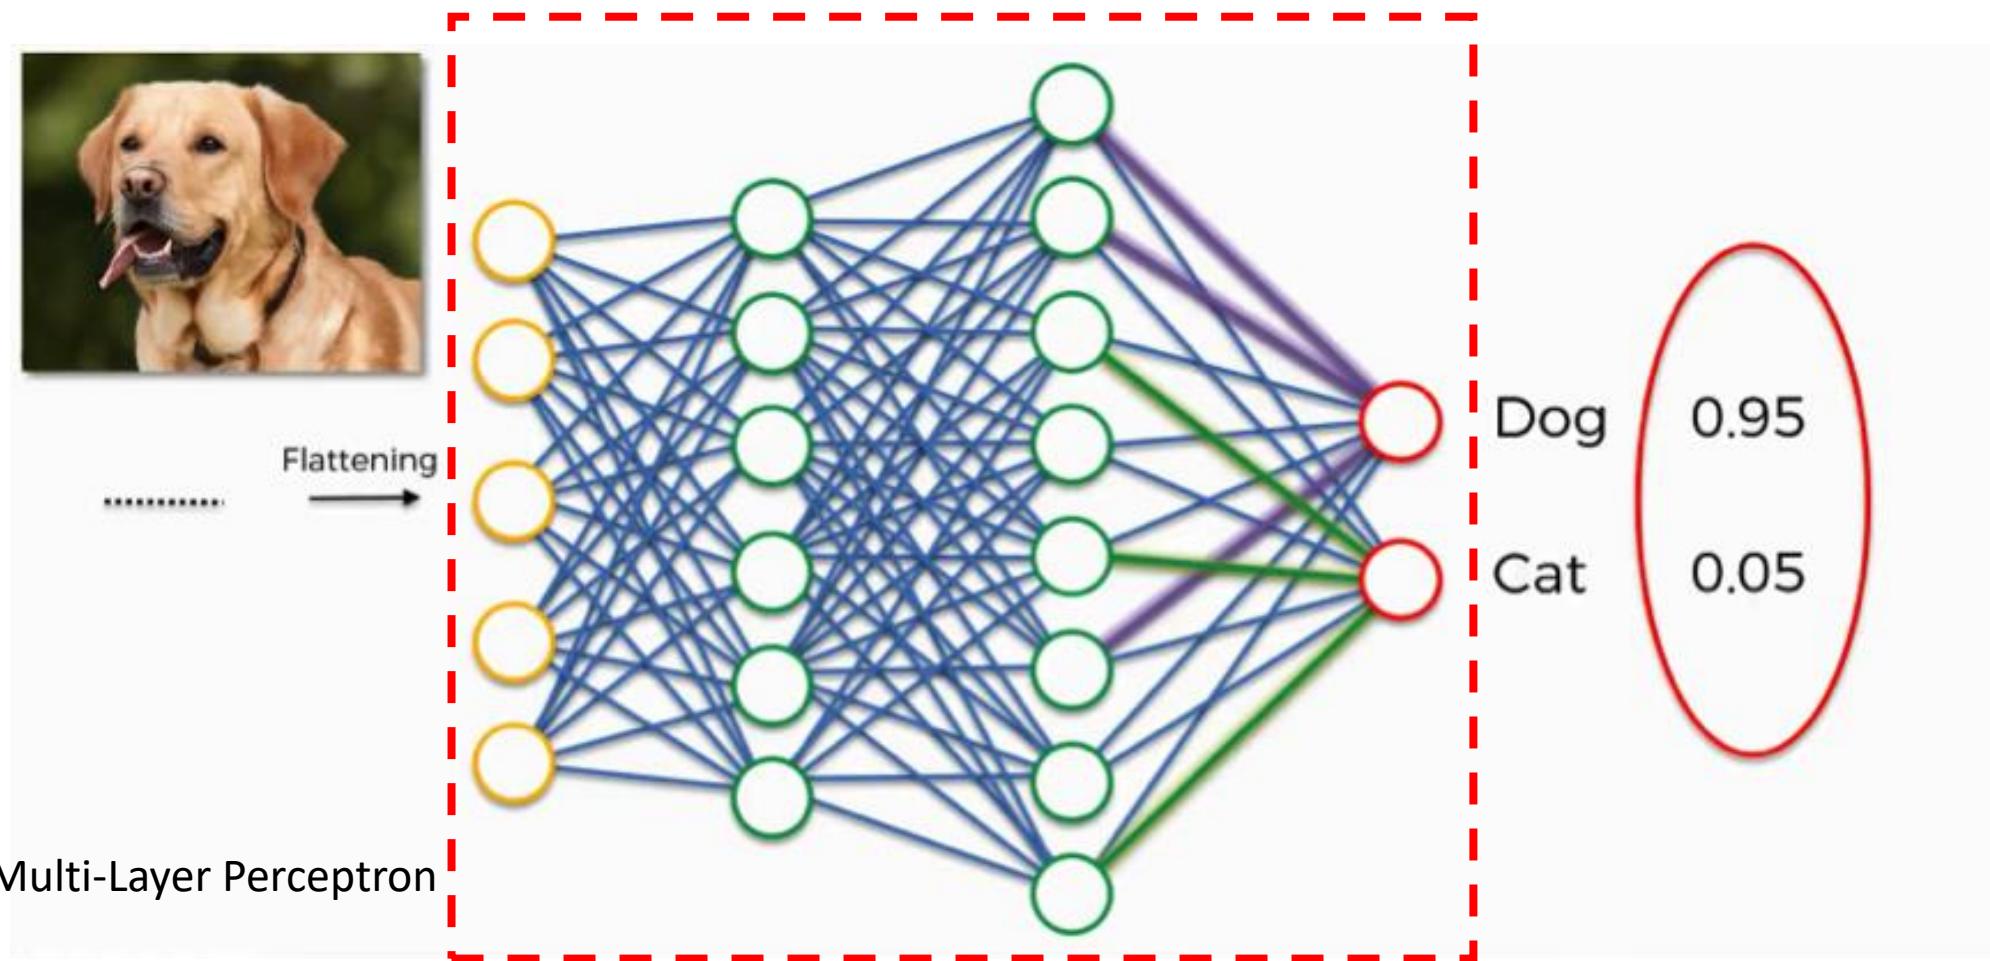
L2 Loss

$$\ell(\hat{y}, y) = (\hat{y} - y)^2$$





Loss – Classification – Cross Entropy Loss

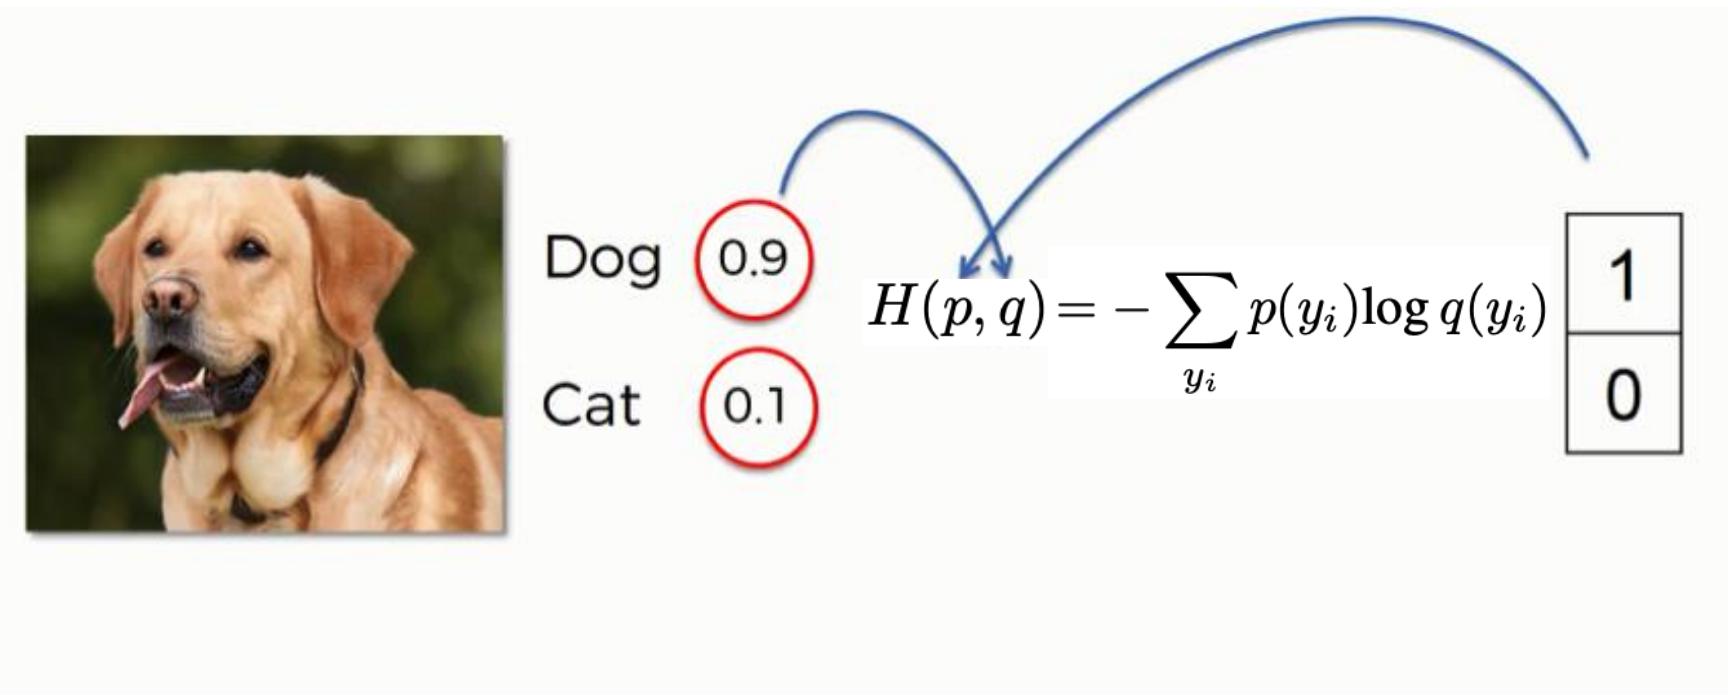




◆ Cross Entropy Loss

$$H(p, q) = - \sum_{y_i} p(y_i) \log q(y_i), \quad \sum_{y_i} p(y_i) = 1, \sum_{y_i} q(y_i) = 1$$

- ◆ $p(y_i = 1)$ is the **ground-truth** probability of category i
- ◆ $q(y_i = 1)$ is the **predicted** probability of category i





Cross Entropy Loss → Negative Log Softmax

$$H(p, q) = - \sum_{y_i} p(y_i) \log q(y_i) \iff L = H(p, q) = -\log \frac{e^{s_{i^*}}}{\sum_j e^{s_j}}, \text{ where } p(y_{i^*}) = 1$$



Unnormalized probabilities						
dog	3.2		24.5		0.13	$L = -\log \frac{e^{3.2}}{e^{3.2} + e^{5.1} + e^{-1.7}}$
cat	5.1	exp	164.0	normalize	0.87	$= -\log \frac{24.5}{24.5 + 164.0 + 0.18}$
frog	-1.7		0.18		0.00	$= -\log 0.13$
						$= 0.89$

Network output **score**:
Unnormalized log probabilities

Probabilities



Multi-Layer Perceptron (MLP)

- How many Perceptrons?

$$4 + 2 = 6$$

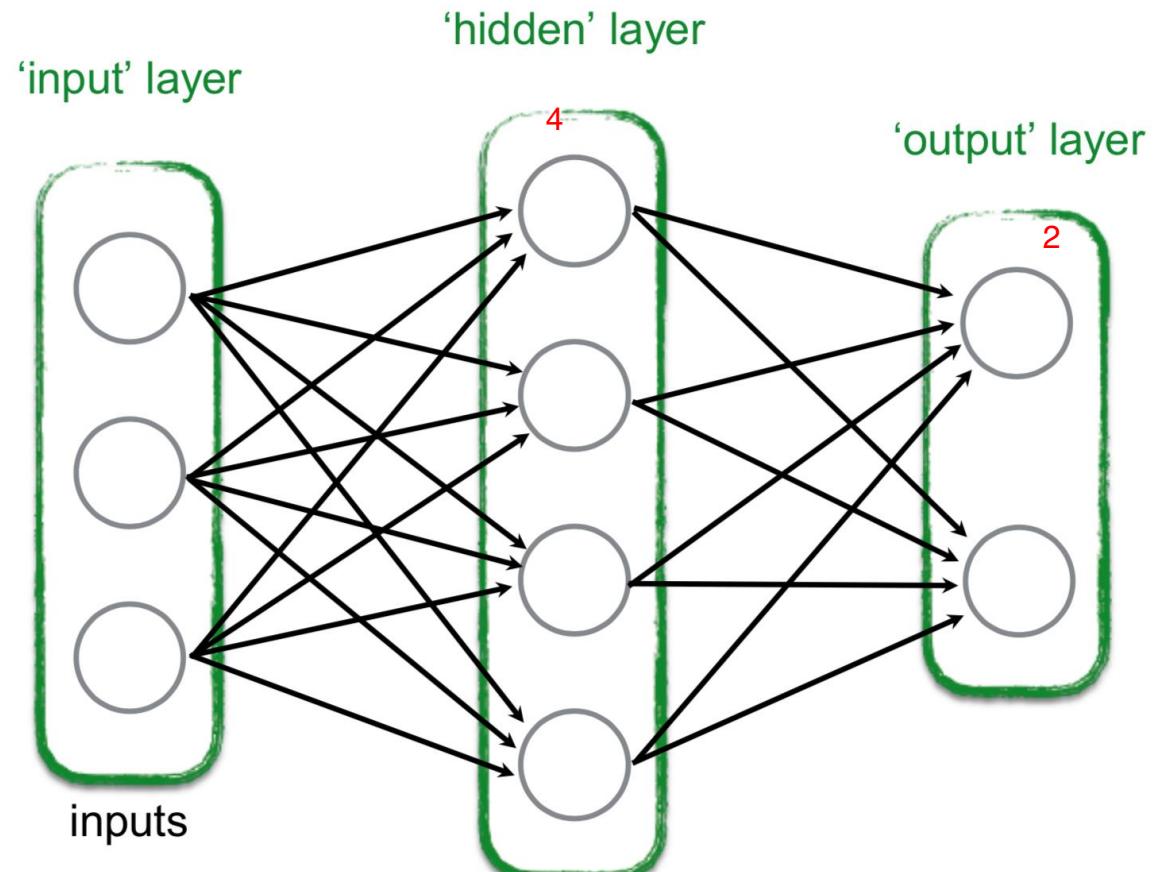
- How many trainable parameters?

$$3 \times 4 + 4 \times 2 = 20$$

- Output $y = w_2^T (w_1^T x) = w'^T x$

- It is **linear** no matter how many layers are in the MLP!

如果没有激活函数





MLP – XOR Problem

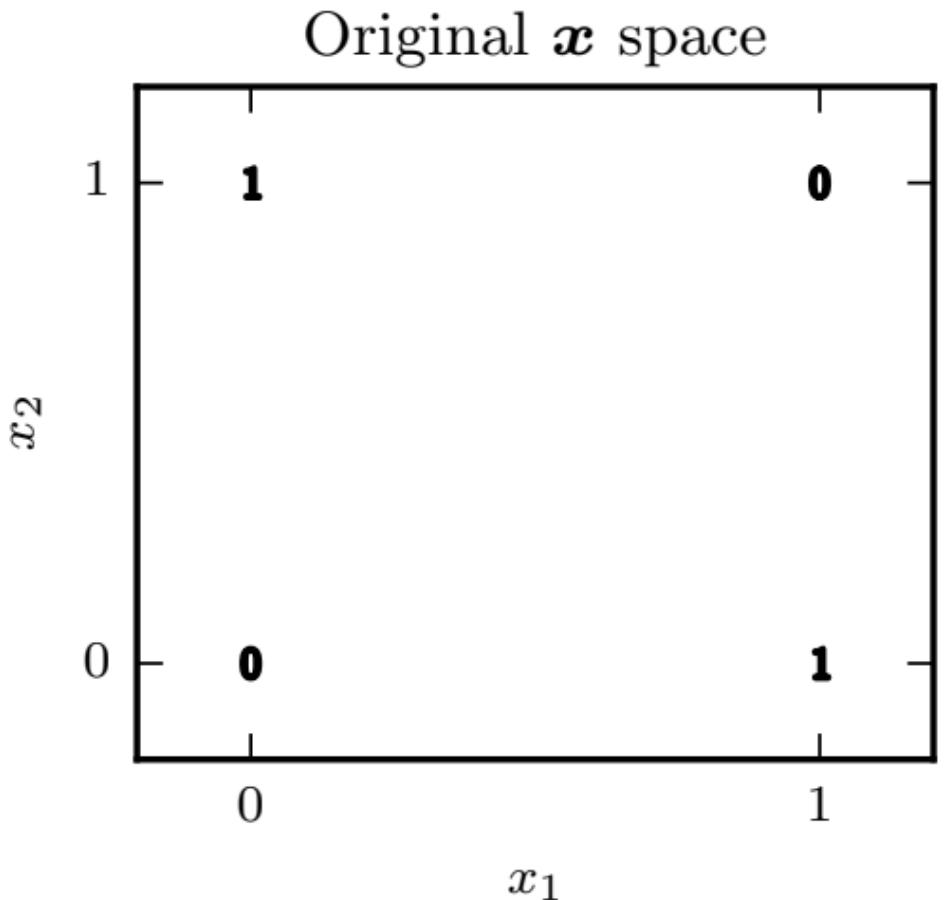
$$\begin{aligned}XOR(0,0) &= XOR(1,1) = 0 \\XOR(0,1) &= XOR(1,0) = 1\end{aligned}$$

Train a linear MLP $y = w^T[x_1, x_2]$ so that:

$$\begin{aligned}w^T[0,0] &= w^T[1,1] < 0 \\w^T[0,1] &= w^T[1,0] > 0\end{aligned}$$

NOT Possible, because $w^T x = 0$ is a hyperplane

How? Make it non-linear



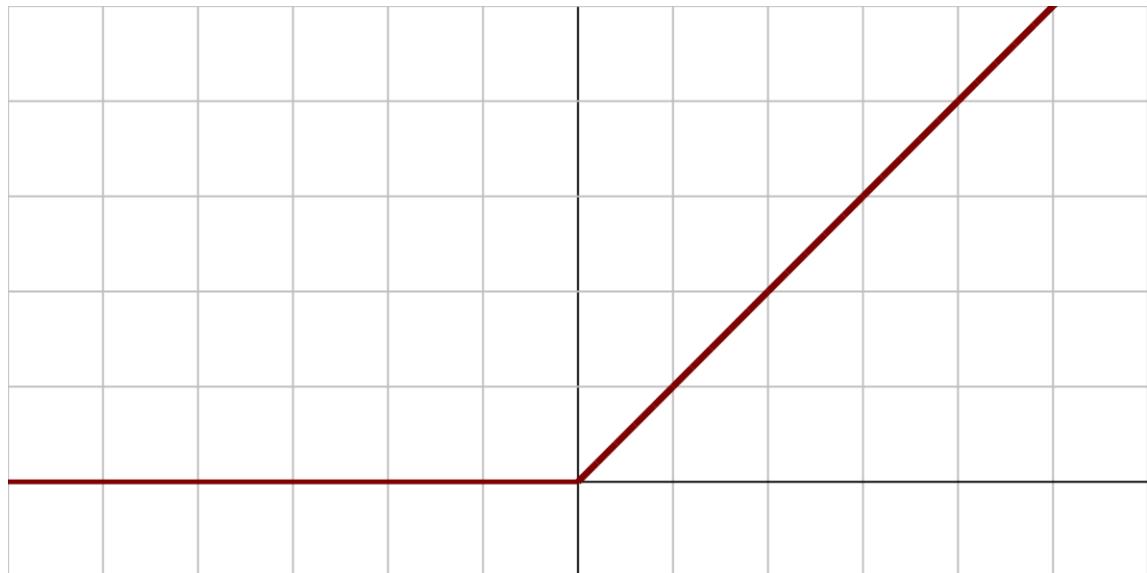
找不到一个平面可以把0和1分开



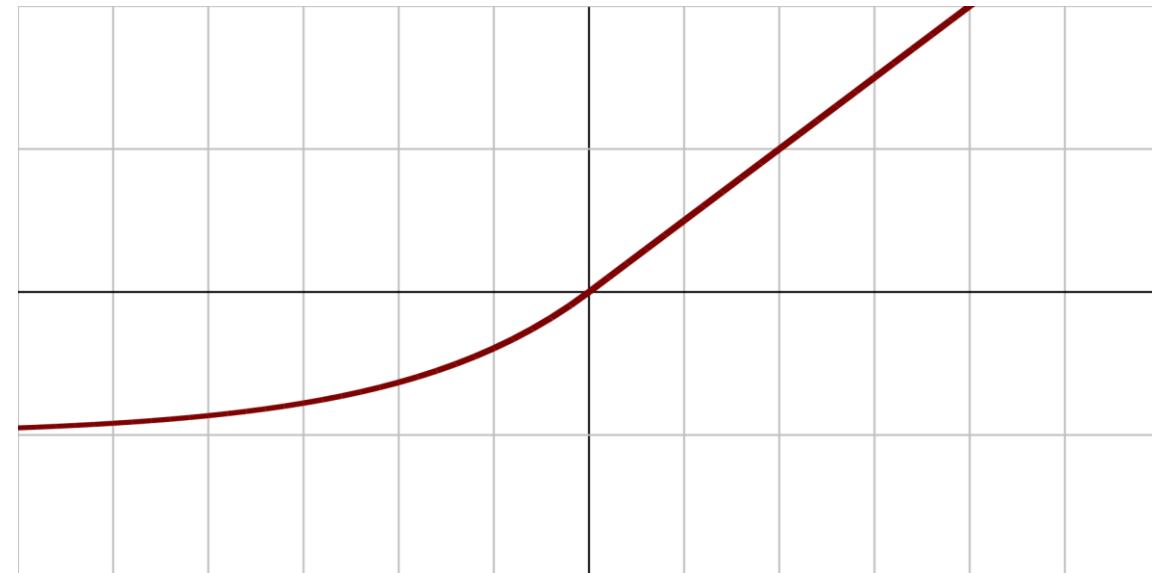
MLP – Activation Function

- Add non-linearity by activation function $g(\cdot)$
- A perception becomes $y = g(f(x)) = g(w^T x)$

Rectified Linear Unit (ReLU) $f(x) = \begin{cases} 0, & \text{for } x \leq 0 \\ x, & \text{for } x > 0 \end{cases}$



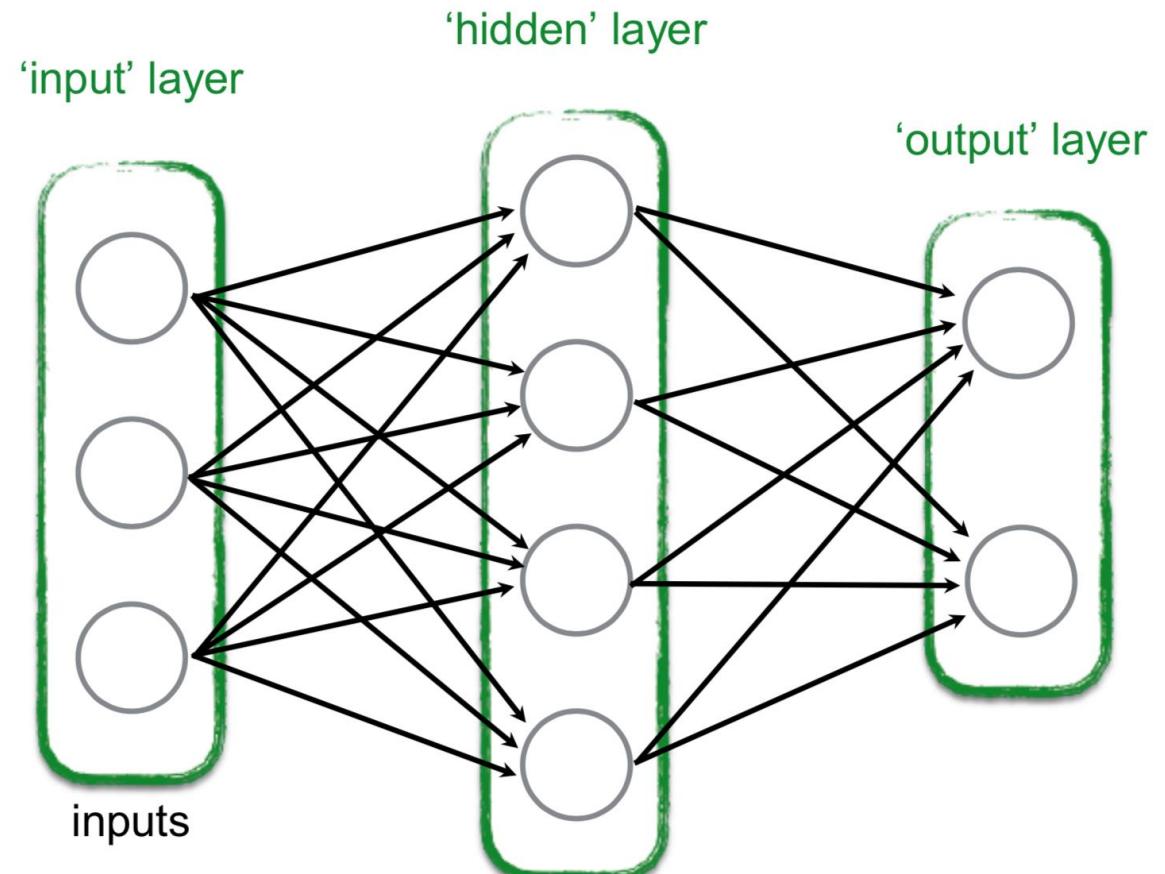
Exponential Linear Unit (ELU) $f(x) = \begin{cases} \alpha(e^x) - 1, & \text{for } x \leq 0 \\ x, & \text{for } x > 0 \end{cases}$





Multi-Layer Perceptron (MLP) Fully Connected Network (FC)

- A MLP with activation & ≥ 1 hidden layers
- Is able to simulate ANY function $f(x)$, where x is the input
- There is proof for this, not shown in this lecture.



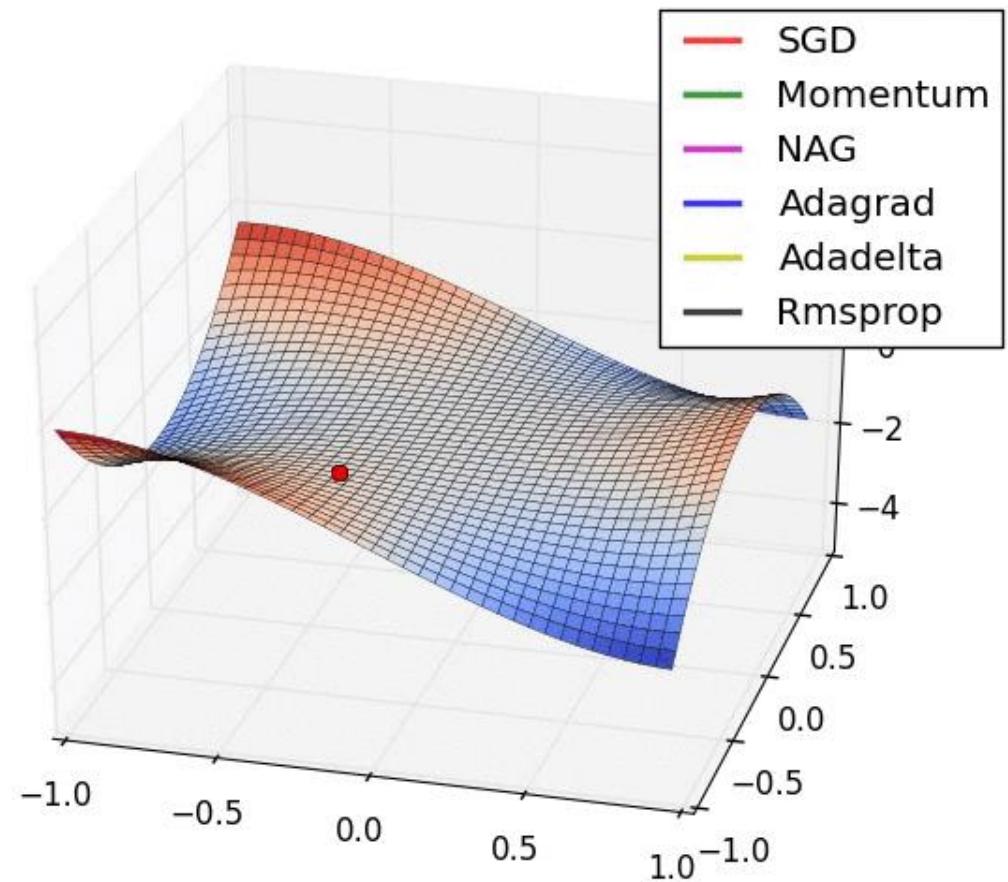


- Same, gradient descent.

$$w_{n+1} = w_n - \lambda \frac{\partial L}{\partial w} \Big|_{w=w_n}, \forall w$$

- What is **Back Propagation**? 更高效的梯度下降
- An efficient way of gradient descent.
- What are optimizers like SGD (Stochastic Gradient Descent), Adam?
- Improvements over gradient descent, still gradient descent.

Gradient Descent Example





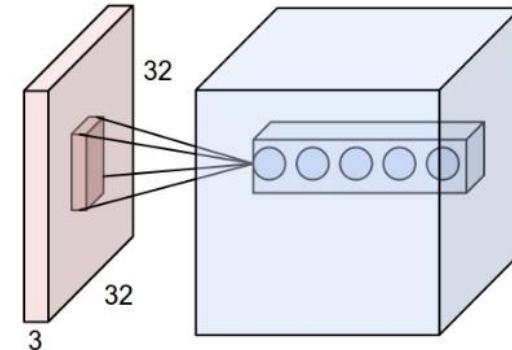
What is CNN?

- Apply convolution on **1D/2D/3D**

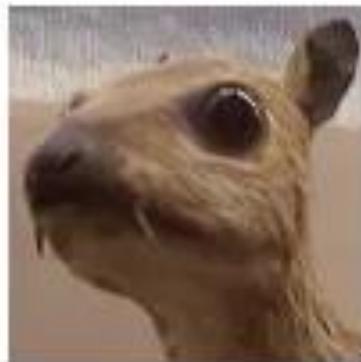
Why CNN?

- Features can be extracted in a **local neighborhood.**

(We will give detailed answers later, after showing what is CNN)



Input image



Convolution
Kernel

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

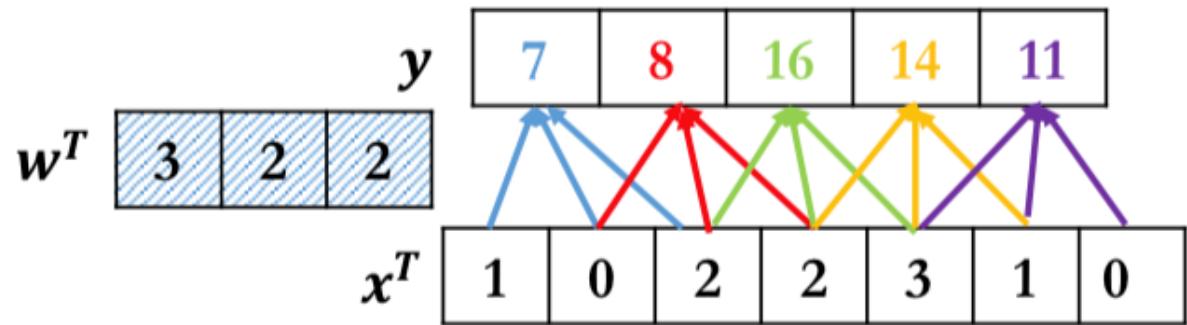
Feature map





1D Convolution

$$y_t = \sum_{i=0}^{k-1} w_i \cdot x_{t+i}$$



• w is the kernel / filter

- Length k
- They are the **unknown/trainable parameters**

• x is the input

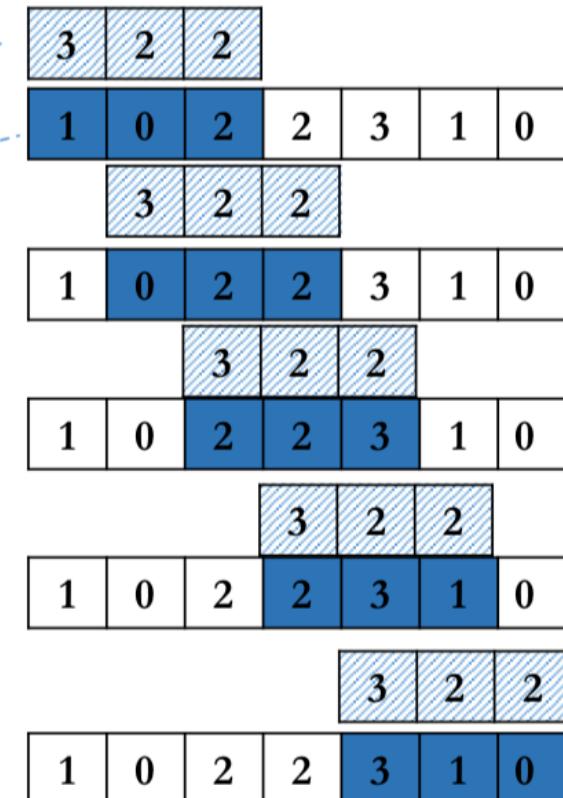
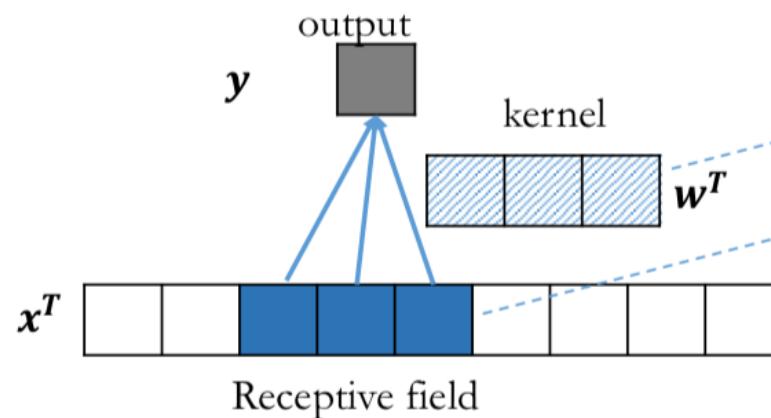
- Length n
- Receptive field: $[x_t, x_{t+1}, \dots, x_{t+k-1}]$
- Each receptive field generates one output value y_t

• y is the output

- Length o



1D Convolution



$$3 \times 1 + 2 \times 0 + 2 \times 2 = 7$$

$$3 \times 0 + 2 \times 2 + 2 \times 2 = 8$$

$$3 \times 2 + 2 \times 2 + 2 \times 3 = 16$$

$$3 \times 2 + 2 \times 2 + 2 \times 1 = 14$$

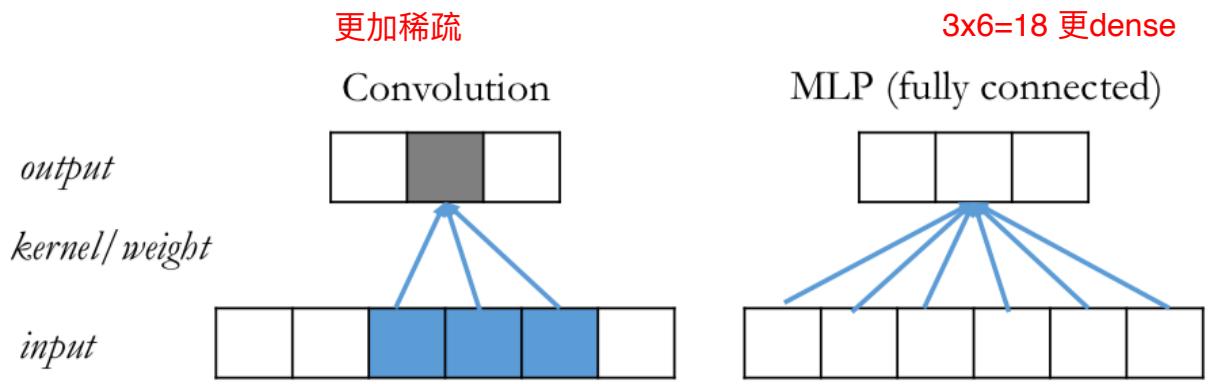
$$3 \times 3 + 2 \times 1 + 2 \times 0 = 11$$



Why CNN ?

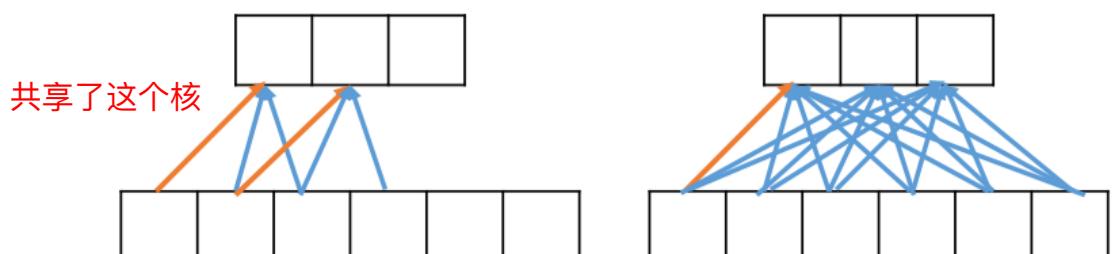
1. Sparse connection vs. Dense

- Fewer parameters. 3 vs. 18
- Less overfitting



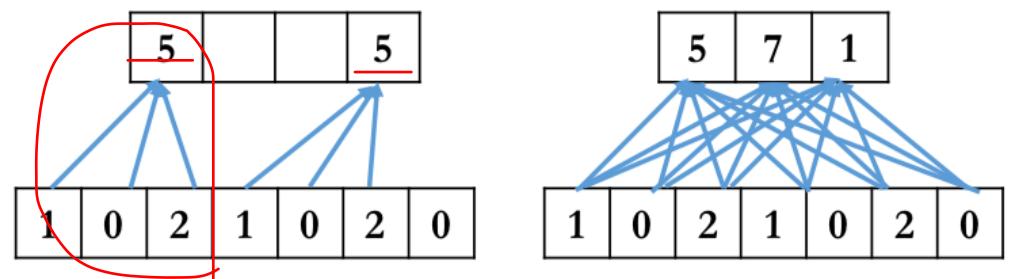
2. Weight sharing vs. Unique weights

- Less overfitting



3. Local invariant vs. Local variant

- Features should not depend on the location within the image
- Make the same prediction no matter where is the object in the image





1D Convolution - Padding

❖ How to determine edge values?

❖ **Padding!**

❖ Given padding p , what is the output length o ?

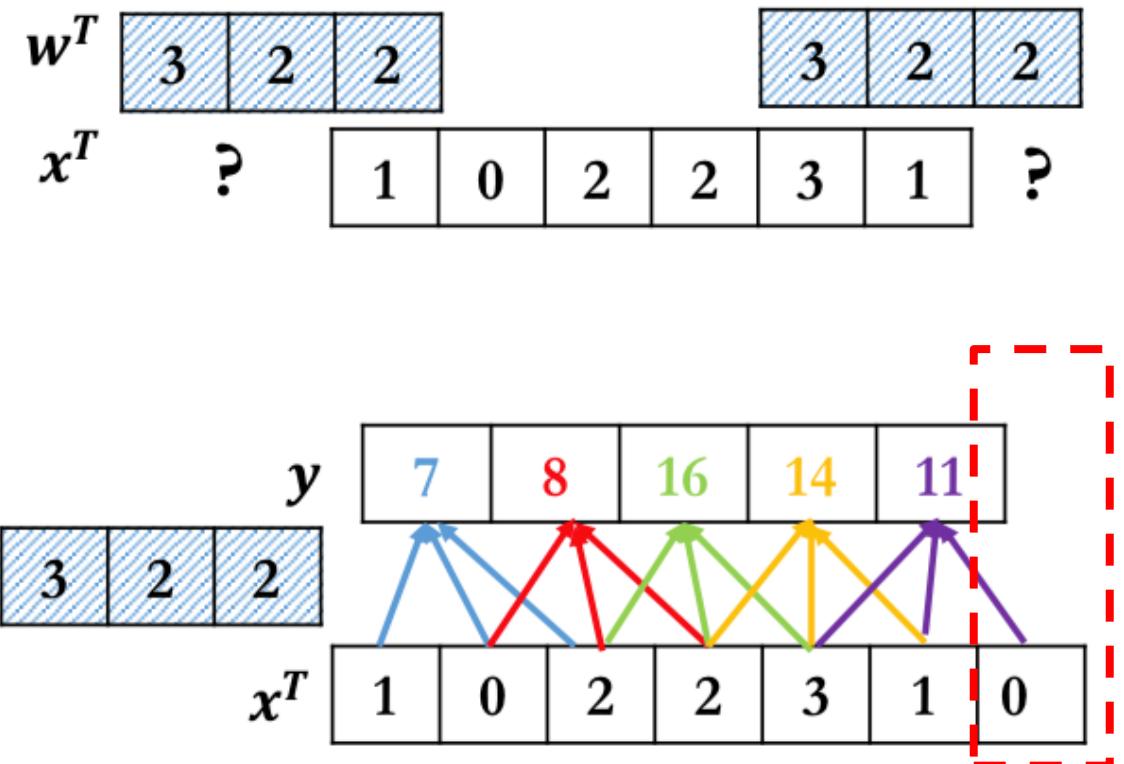
- kernel size k
- Input length n

❖ Without padding:

$$o = n - k + 1$$

❖ With padding:

$$o = (n + p) - k + 1$$





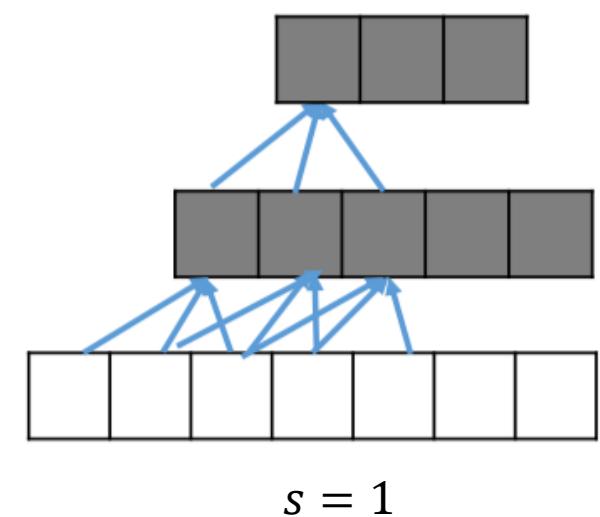
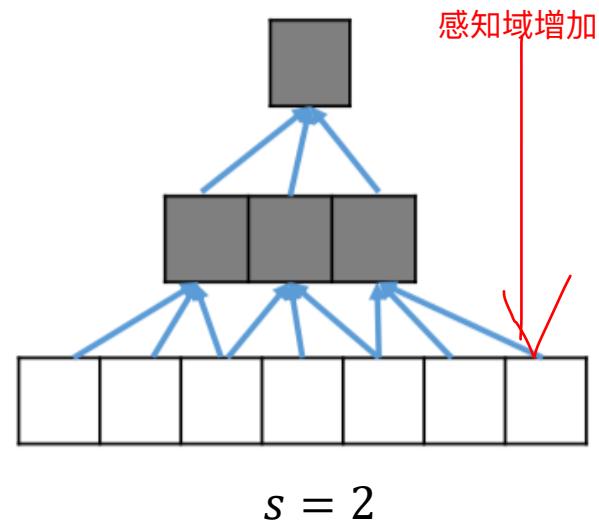
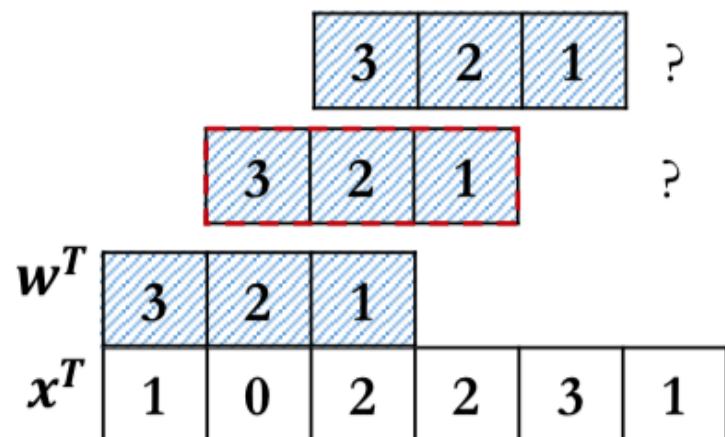
1D Convolution - Stride

• How many steps to move for the next receptive field?

- Stride! Skip some elements when $s > 1$

• Why?

- Less compute
- Increase receptive field

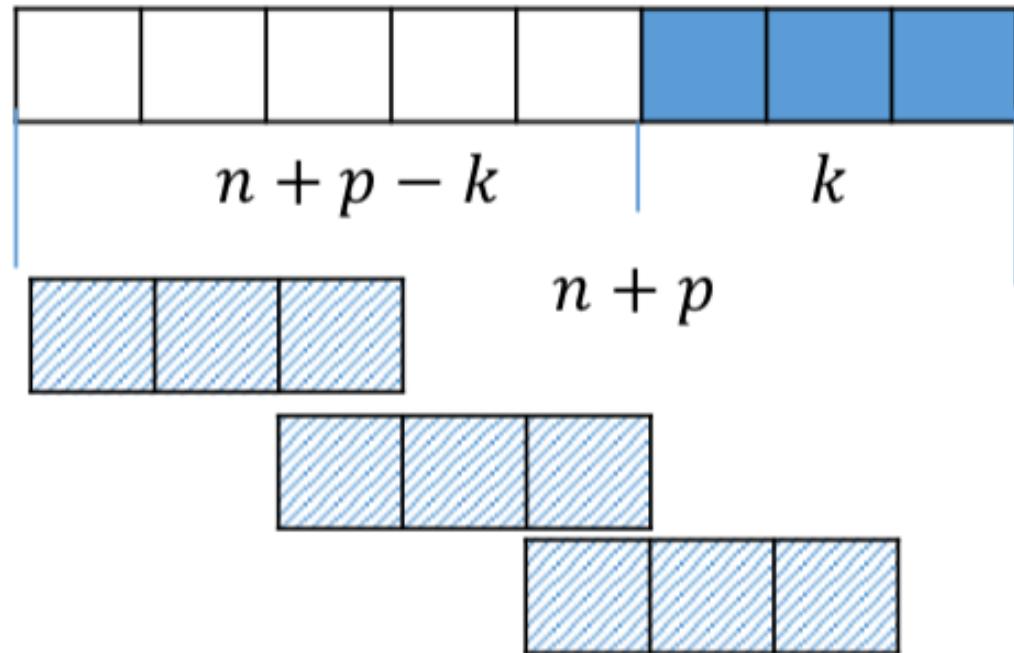




What is the output length o ?

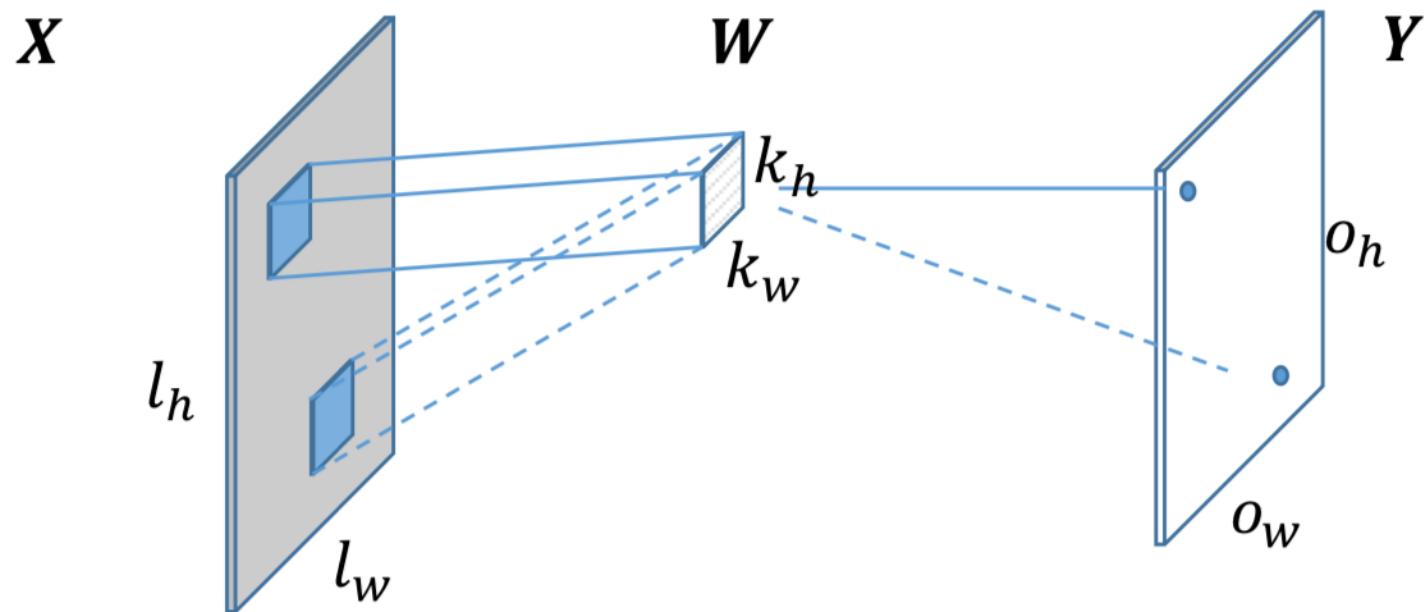
- kernel size k
- Input length n
- Padding p
- Stride s

$$\bullet \quad o = \left\lfloor \frac{n+p-k}{s} \right\rfloor + 1$$





Same but everything 2D

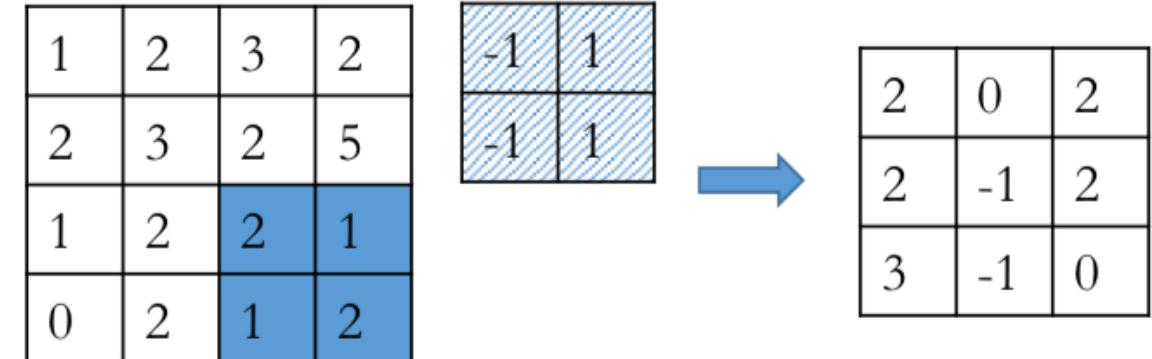
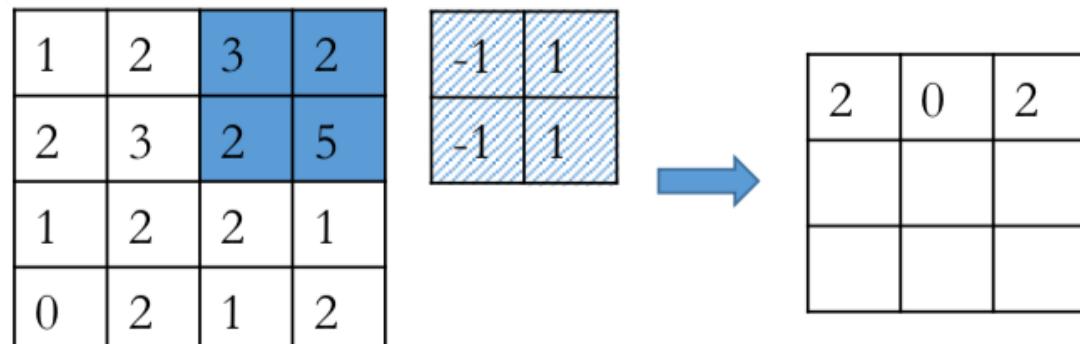
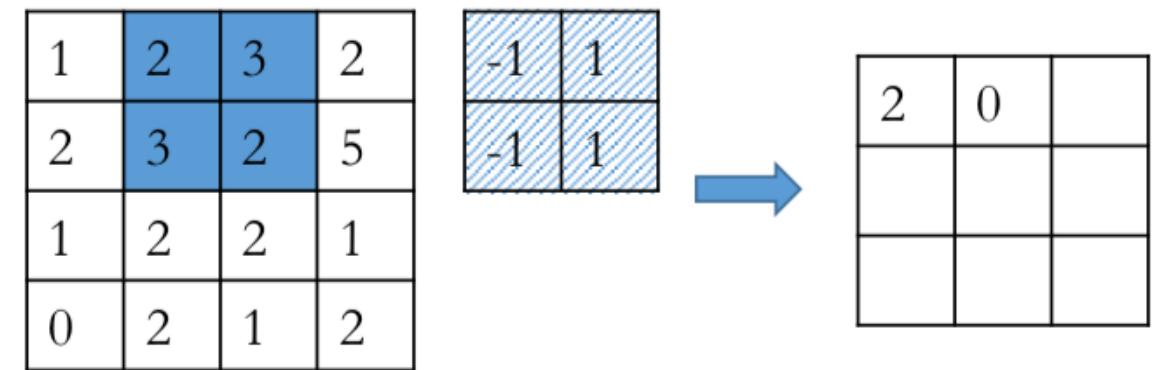
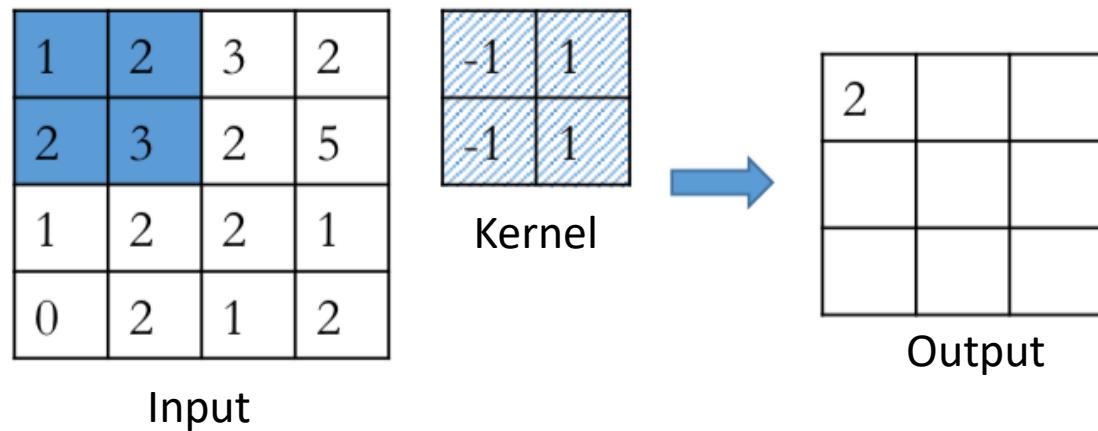


$$Y_{i,j} = \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} X_{i*s_h+a, j*s_w+b} \times W_{a,b}$$



2D Convolution

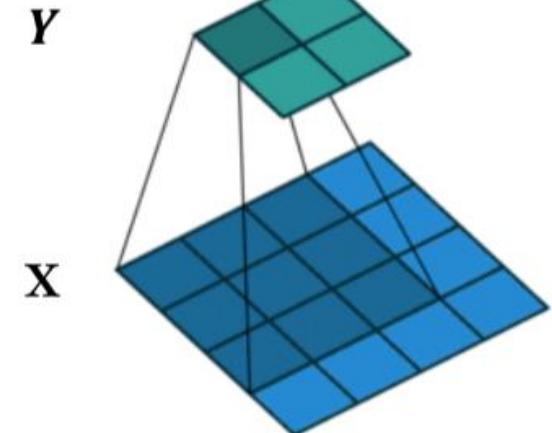
$$Y_{i,j} = \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} X_{i*s_h+a, j*s_w+b} \times W_{a,b}$$



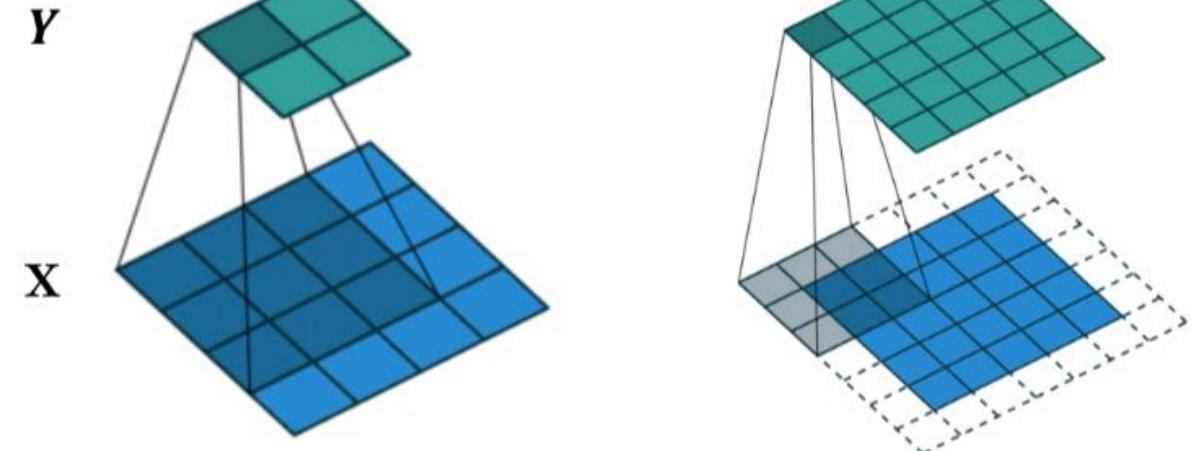


2D Convolution

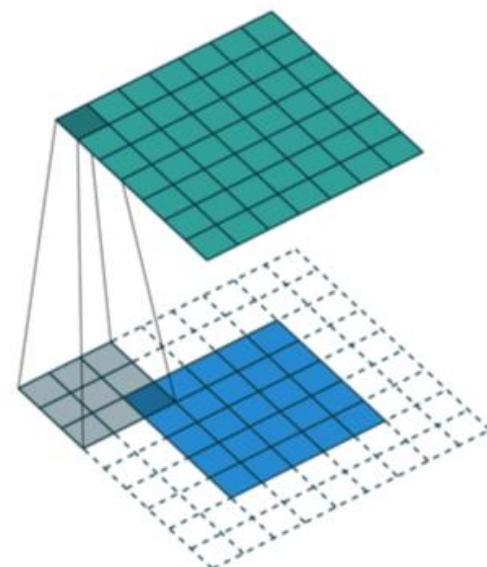
kernel size k
Padding p
Stride s



$k=3, p=0, s=1$ (Valid)



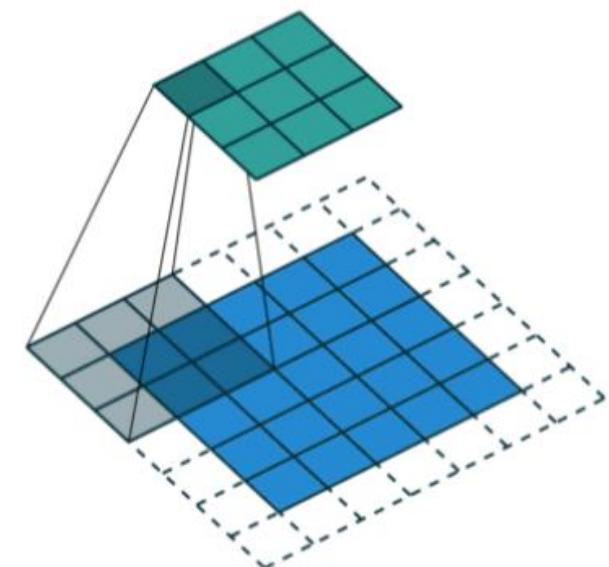
$k=3, p=2, s=1$ (Same)



$k=3, p=4, s=1$ (Full)

Valid: No padding

Same: Add padding so
input size = output size

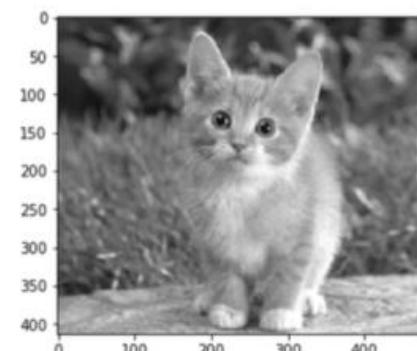
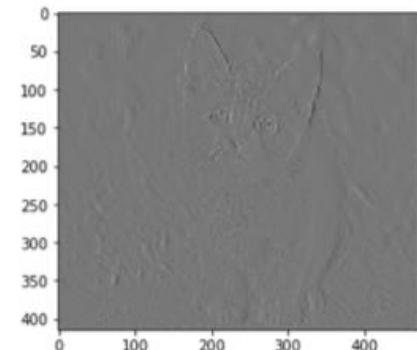
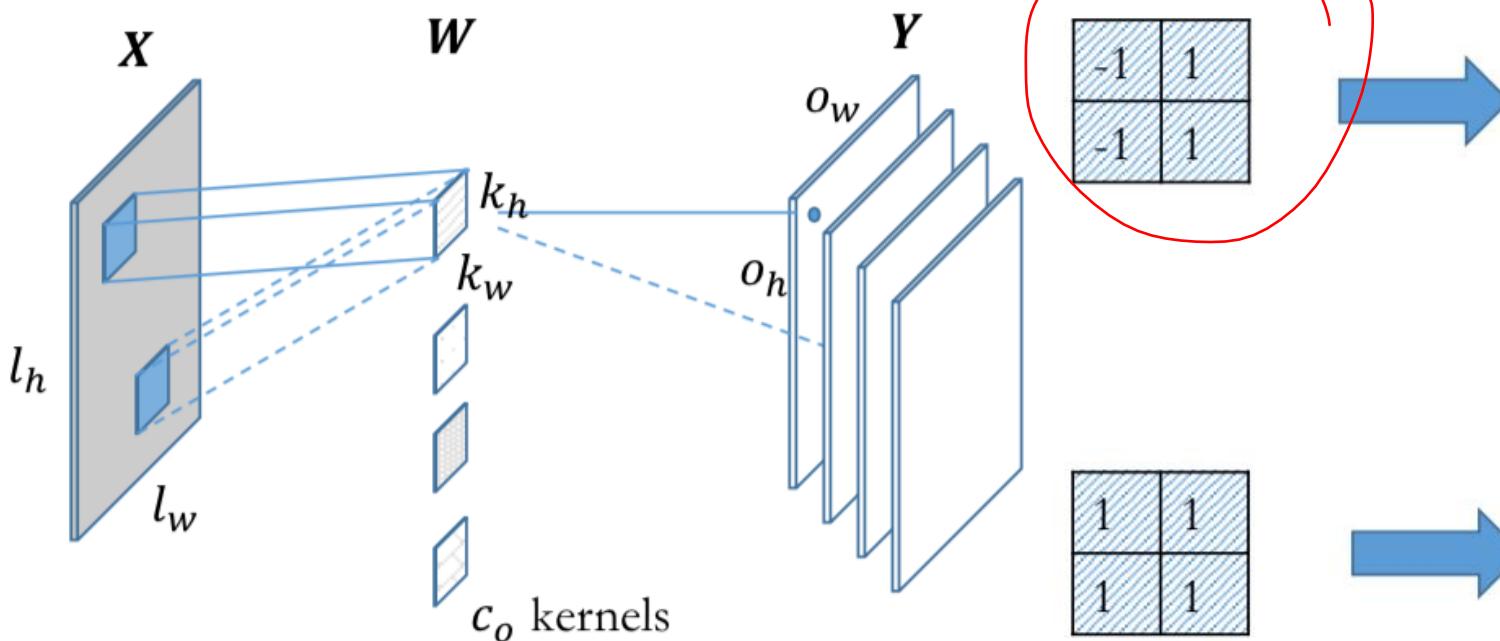
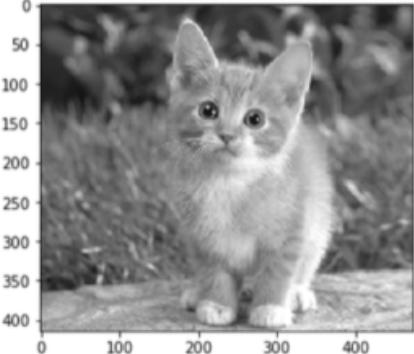


$k=3, p=2, s=2$



2D Convolution – Multiple kernels

◆ Multiple features → Multiple kernels



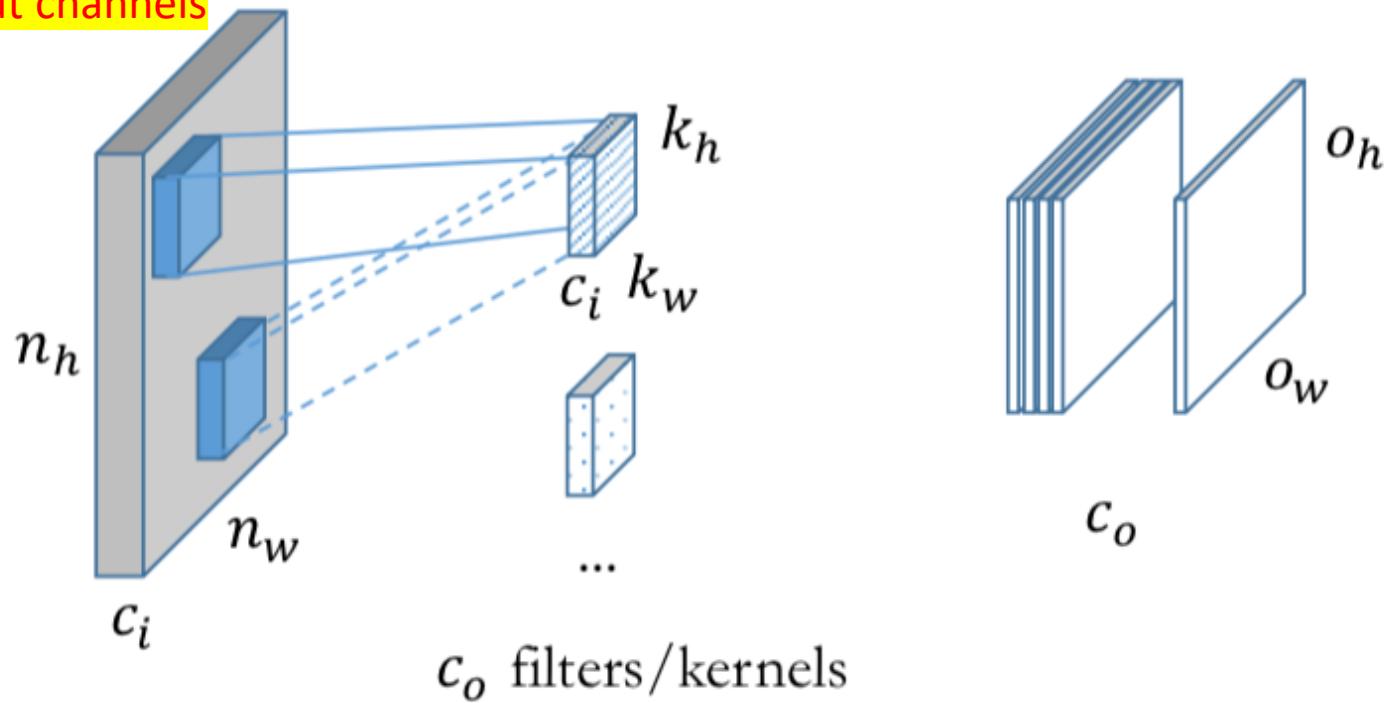


2D Convolution – Multiple Inputs & Kernels

$$\bullet Y_{l,i,j} = \sum_{d=0}^{c_i-1} \sum_{a=0}^{k_h-1} \sum_{b=0}^{k_w-1} X_{d,i+a,j+b} \times W_{l,d,a,b} + b_l, l \in [0, c_o)$$

2D convolution across the plane (indexed by a, b)

Each kernel covers all input channels
(indexed by d)

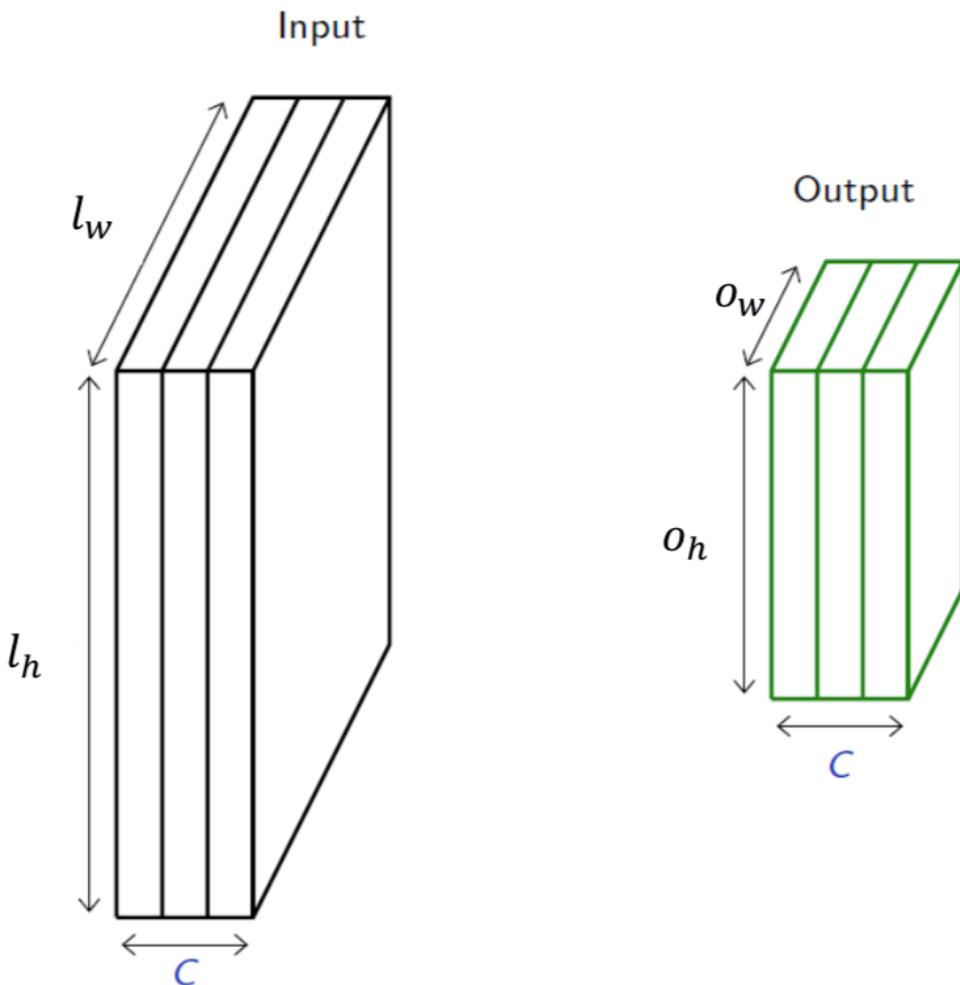




2D Convolution - Pooling

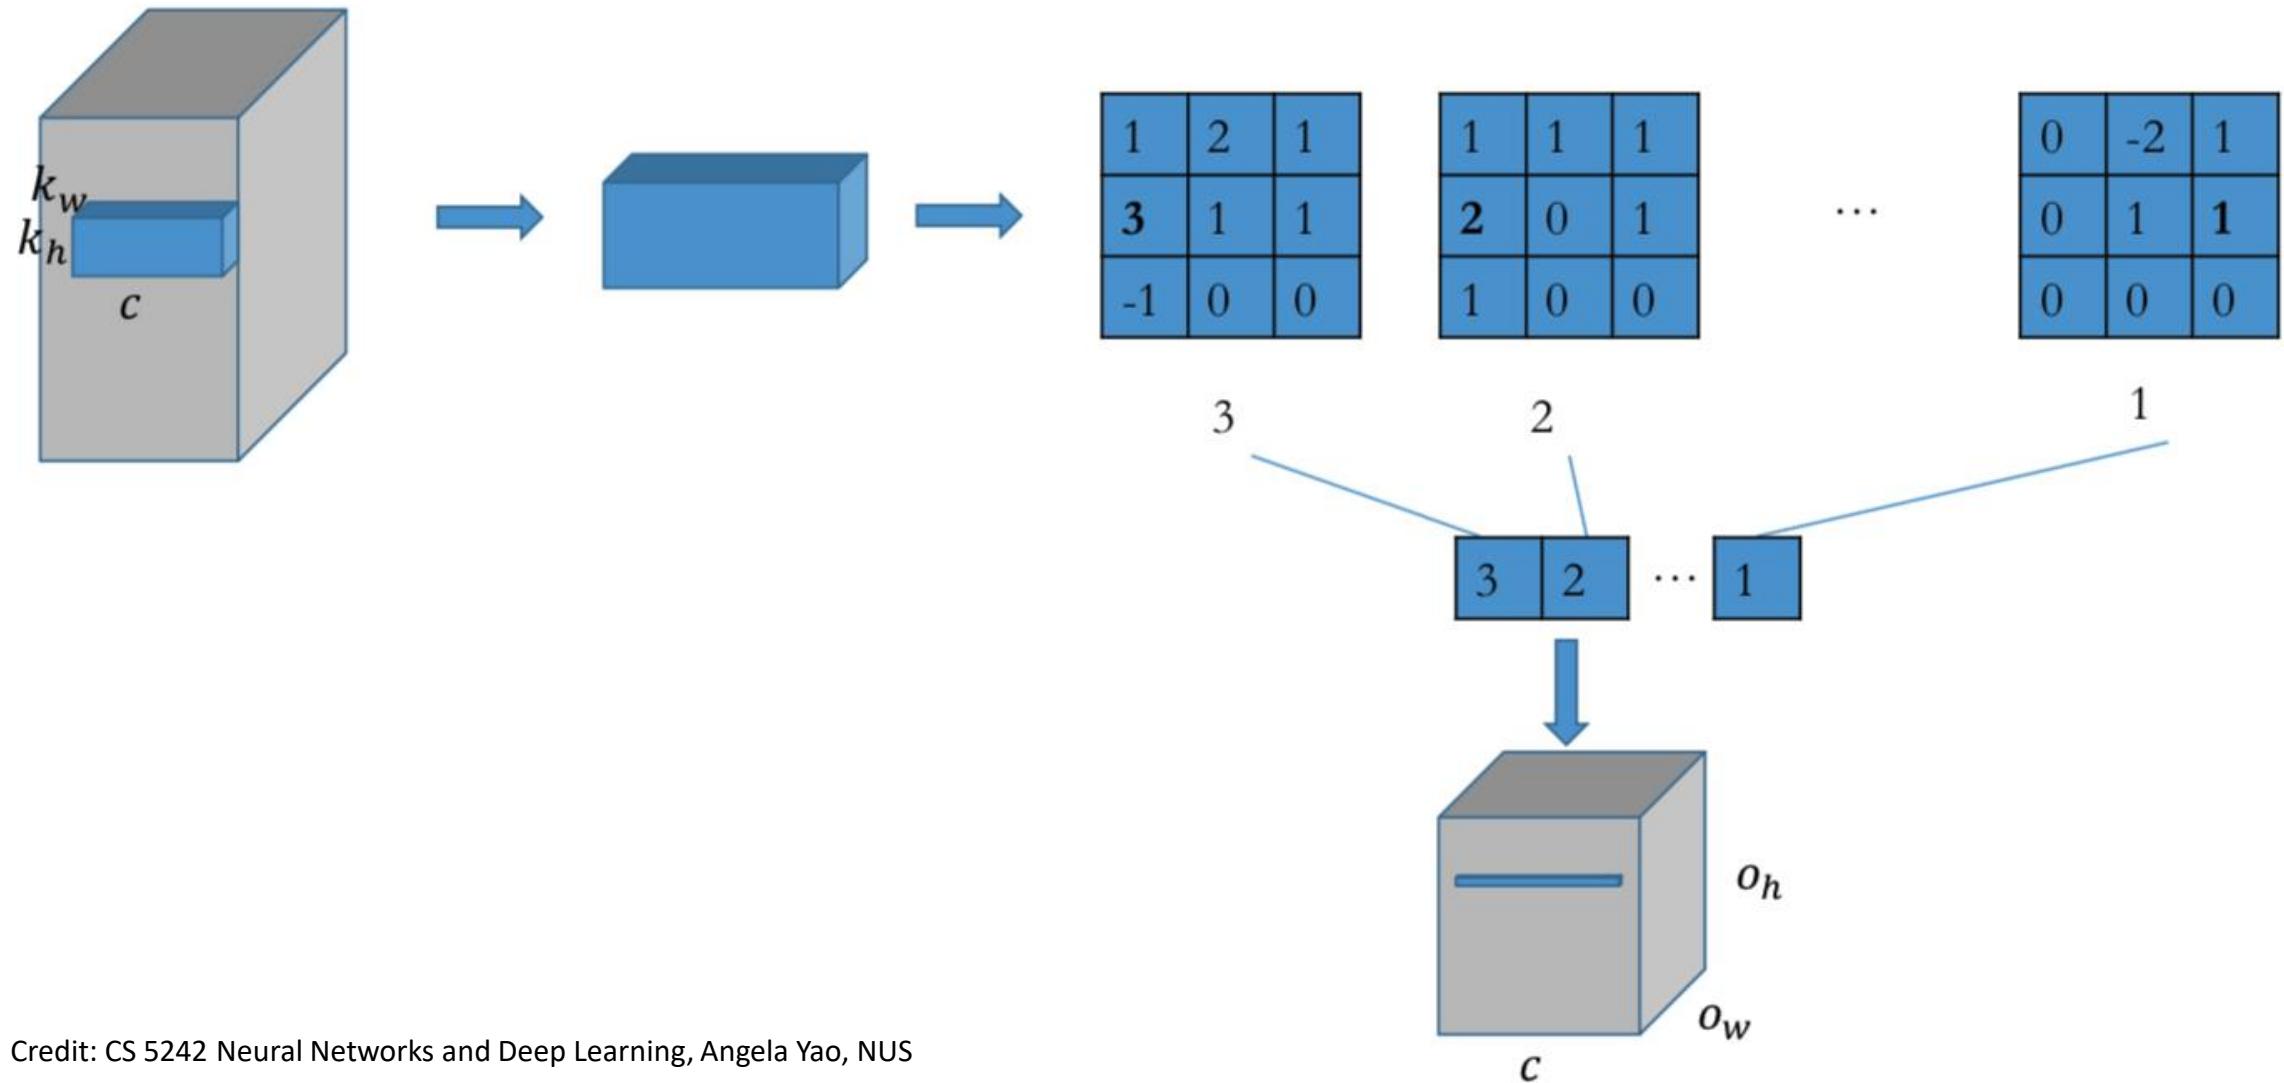
- ◆ Aggregate information in each receptive field
 - Max
 - Average
- ◆ No trainable parameter
- ◆ # input channels = # output channels
 - $c_i = c_o$
- ◆ Same padding/stride method

Pooling其实就是简化的conv





2D Convolution – Max Pooling





- There are c_o kernels, they share same stride and padding, NOT parameter
- Parameter size $c_o \times k_h \times k_w$
- Output shape

- $(c_o, o_h, o_w) = (c_o, \left\lfloor \frac{n_h + p_h - k_h}{s_h} \right\rfloor + 1, \left\lfloor \frac{n_w + p_w - k_w}{s_w} \right\rfloor + 1)$

- Computation cost

- $O((c_o \times k_h \times k_w) \times (o_h \times o_w))$

- **1D Convolution Multiple inputs/kernels**

- Same as 2D
- Set k_h or k_w to be 1.



3D Convolution

◆ Natural extension of 2D Convolution

◆ Blue: input

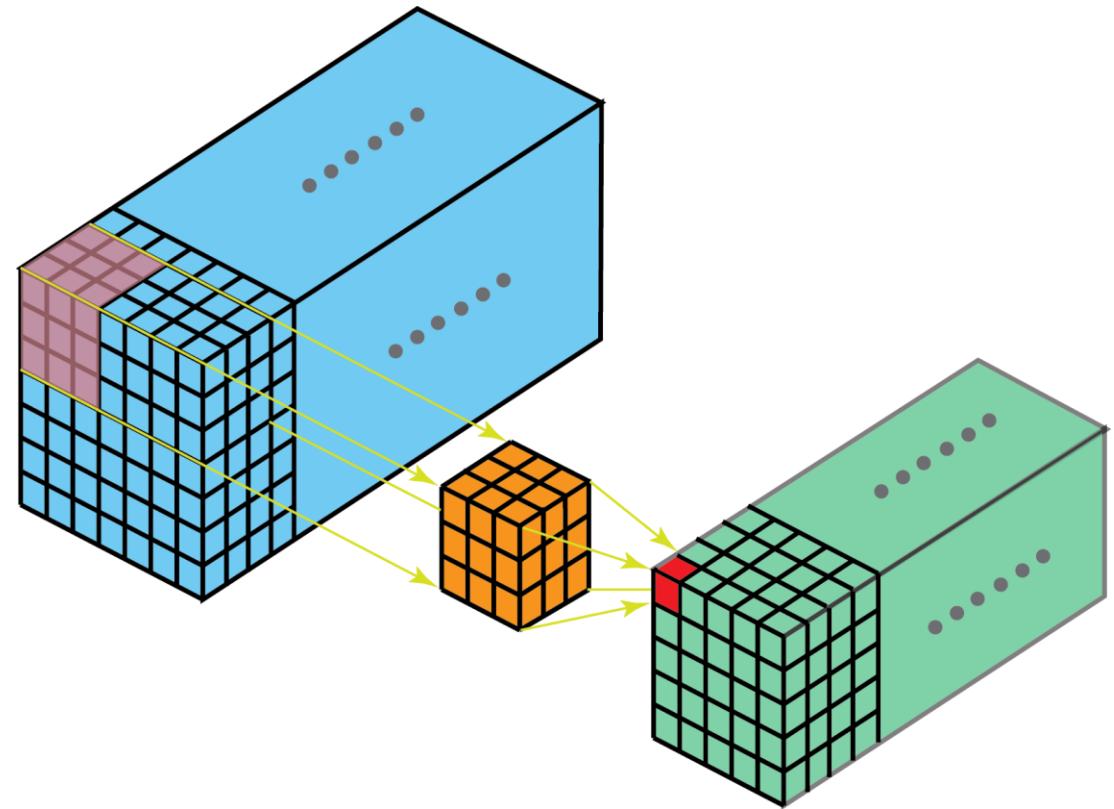
- Each small cube contains d features/channels

◆ Orange: kernel

- Each small cube contains d weights

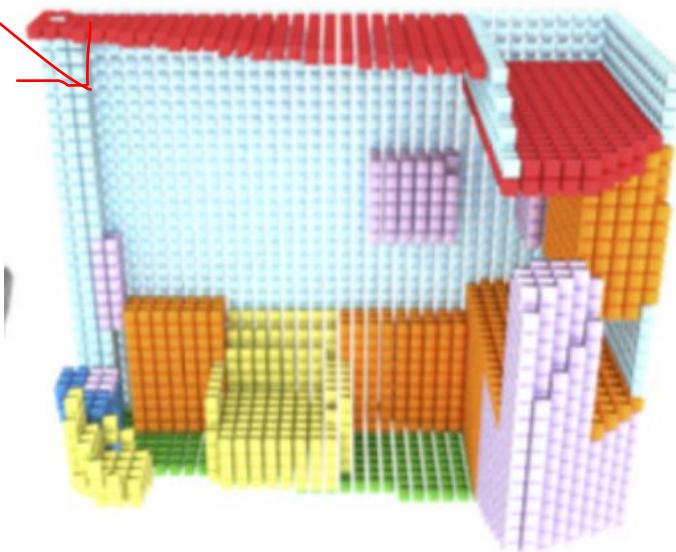
◆ Green: output

- Each small cube is a scalar.
- There are o green blocks





- 3D convolution
- Multi-view projection onto images + 2D convolution
- Simply run 1D/2D convolution or even MLP on point cloud



3D shape model
rendered with
different virtual cameras



x1	y1	z1
x2	y2	z2
x3	y3	z3
x4	y4	z4



Content of each grid cell

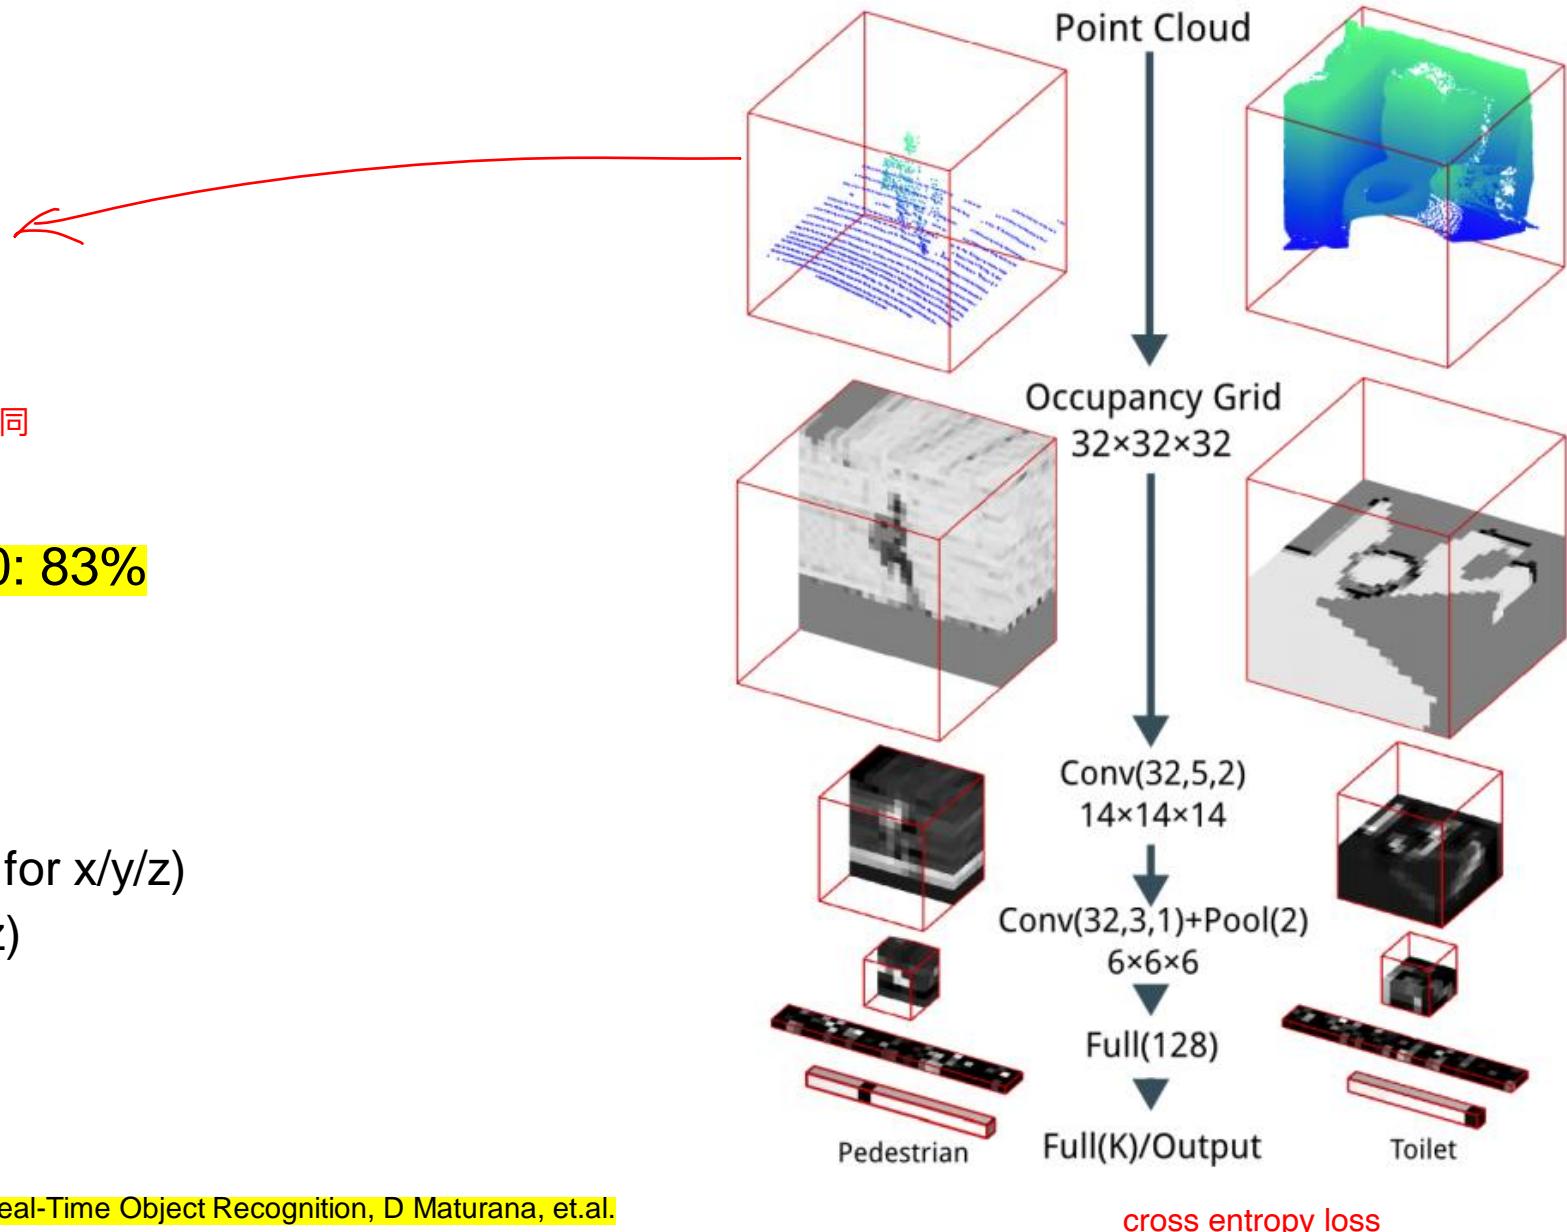
- Binary
- Number of points
- Probability
- etc.

放的东西可以不同

Accuracy on ModelNet40: 83%

Conv(o , k , s)

- o : number of kernels
- k : size of kernel (same for x/y/z)
- s : stride (same for x/y/z)

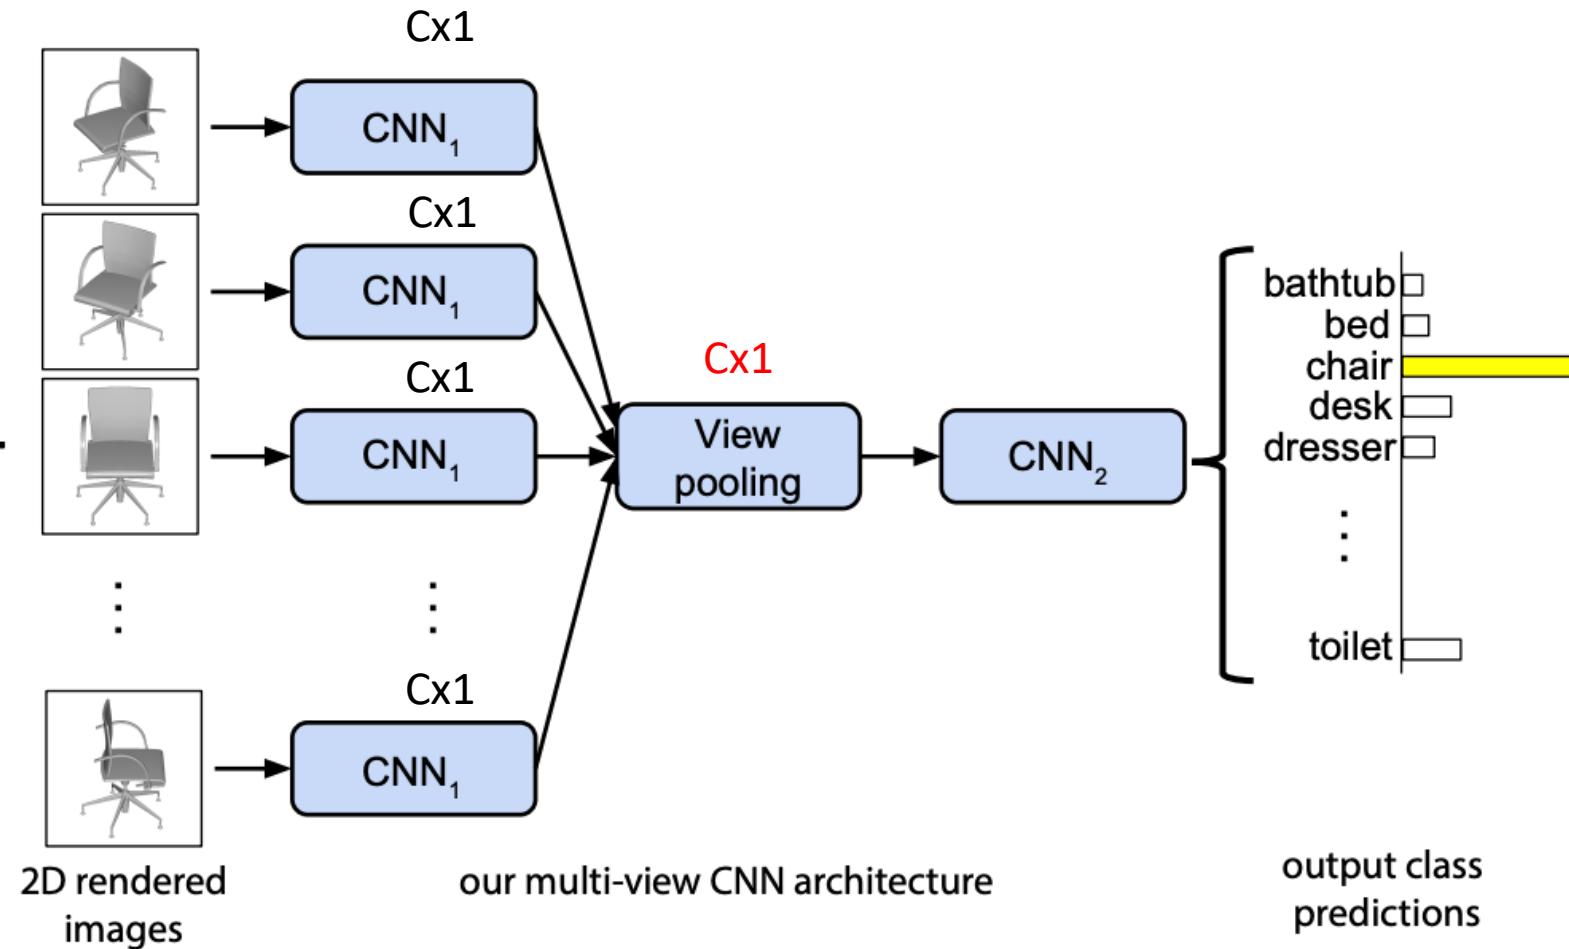




ModelNet40 Accuracy 90.1%



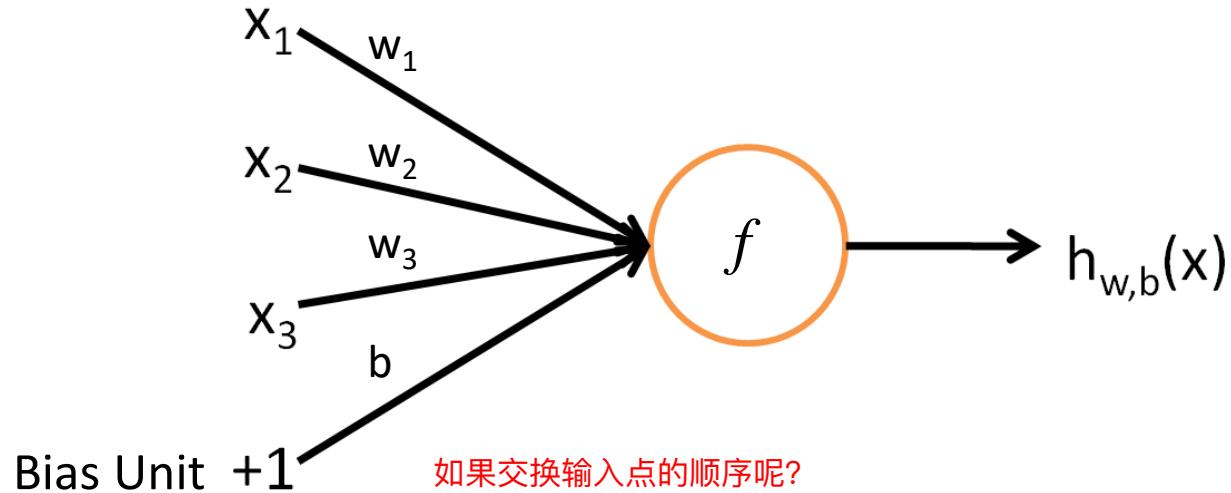
3D shape model
rendered with
different virtual cameras



投影多少次呢



MLP for Point Cloud?

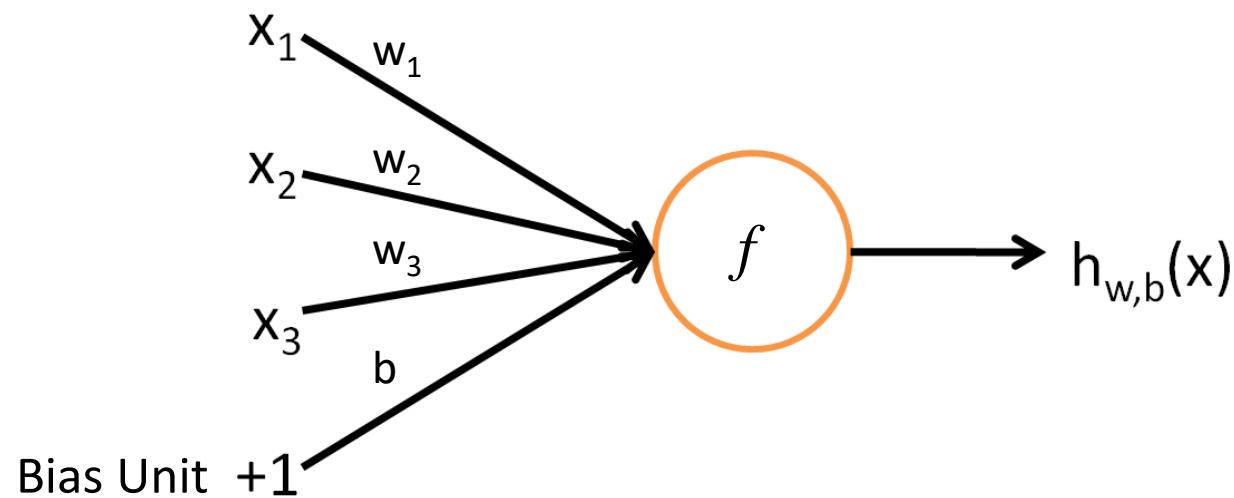


Activation function

$$\begin{aligned} h_{W,b}(x) &= f(W^T x) = f\left(\sum_{i=1}^3 w_i x_i + b\right) \\ &= f(w_1 x_1 + w_2 x_2 + w_3 x_3 + b) \end{aligned}$$



Not Permutation Invariant!

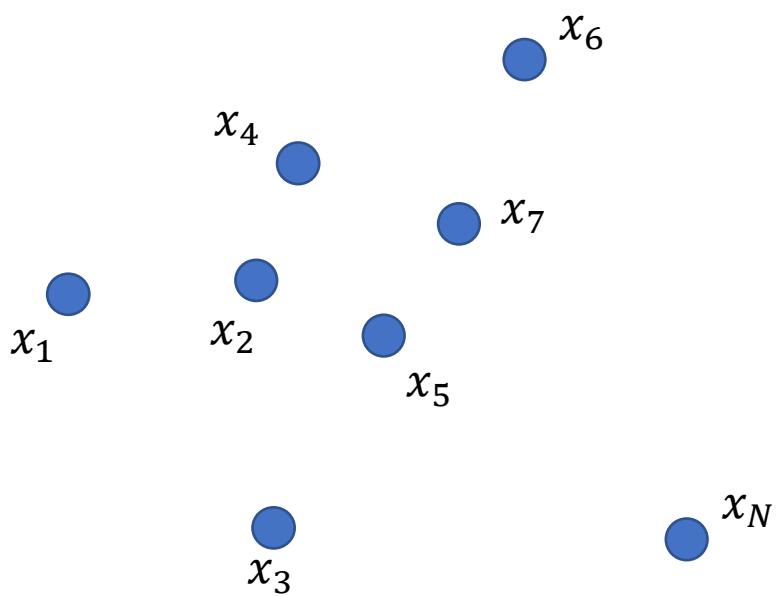


$$f(w_1x_3 + w_2x_1 + w_3x_2 + b) \neq h_{W,b}(x)$$



Enumerate All Permutation?

Point CNN



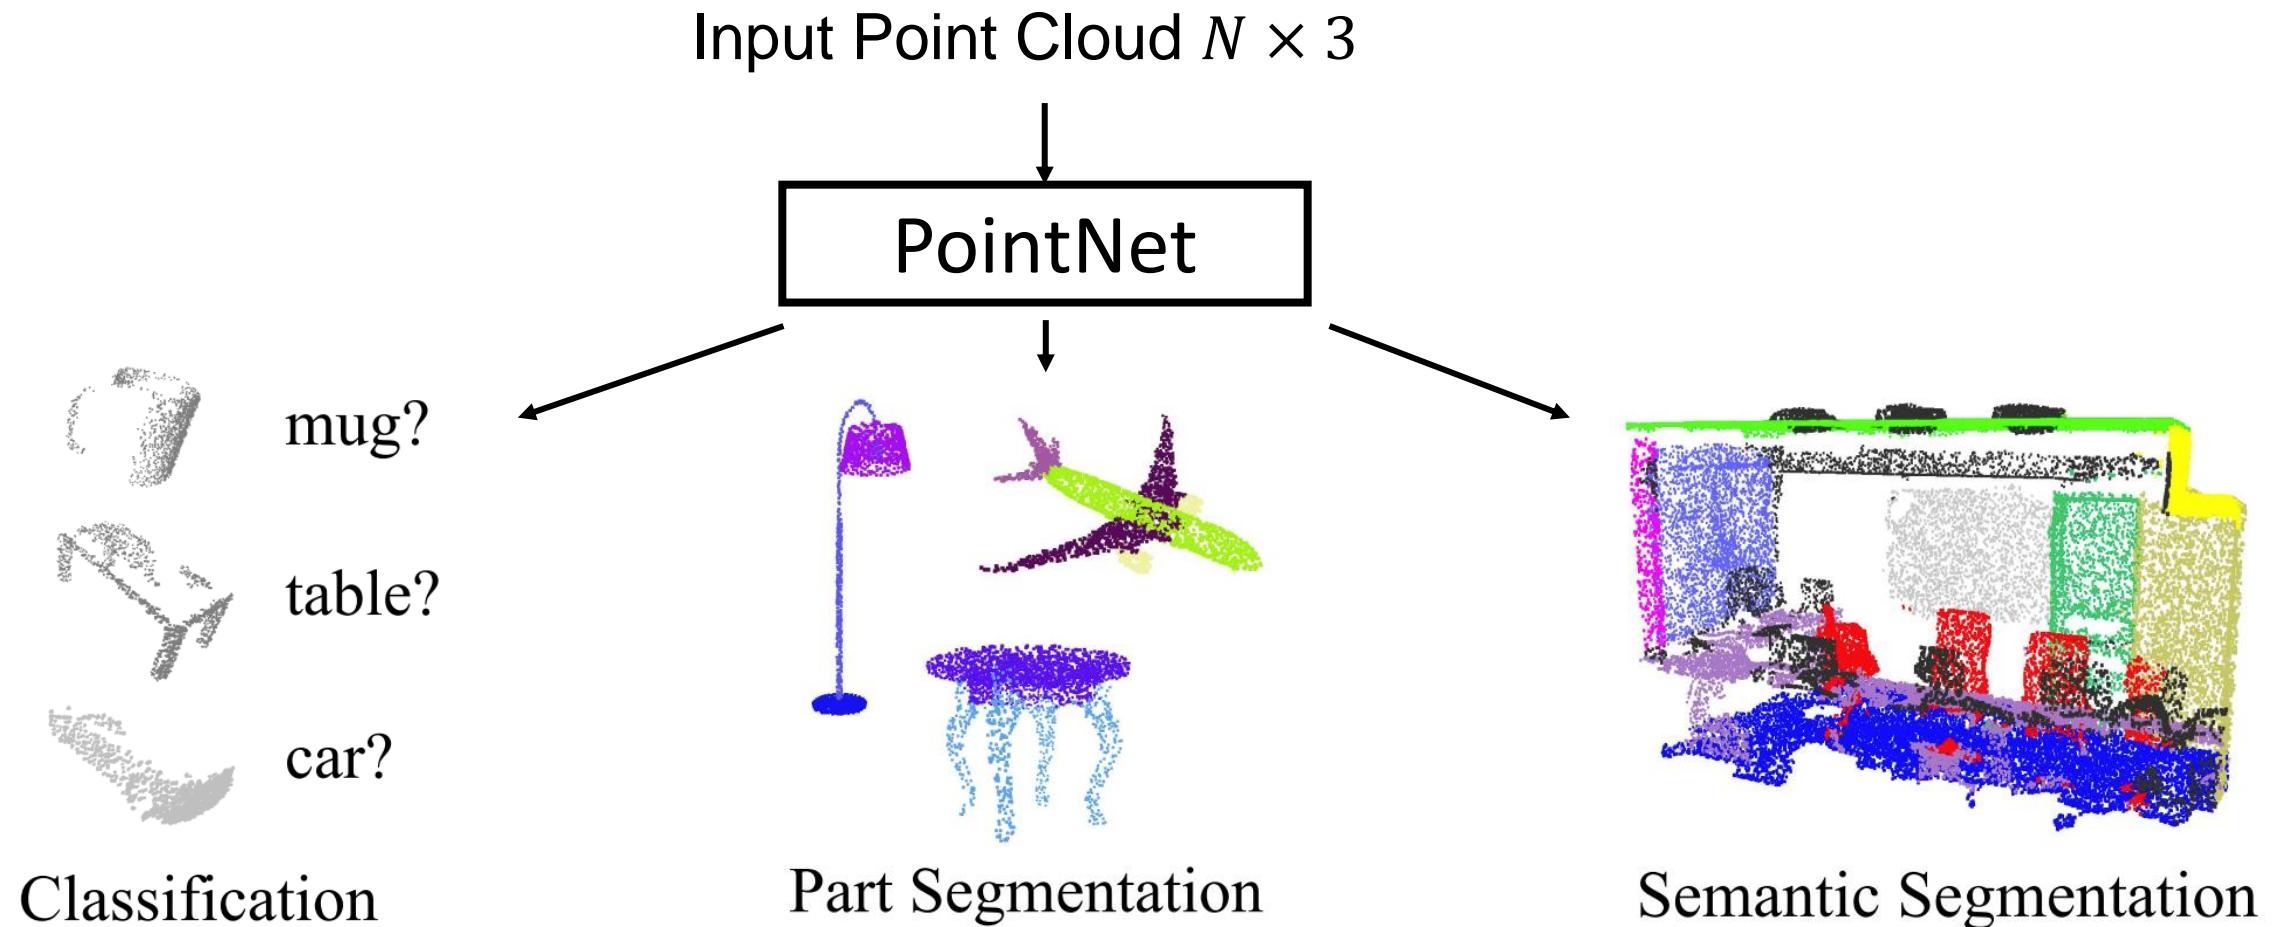
$$X_1 = [x_1^T, x_2^T, \dots, x_N^T]^T$$

$$X_2 = [x_2^T, x_1^T, \dots, x_N^T]^T$$

⋮
⋮

$$X_{N!} = [x_i^T, x_j^T, \dots, x_n^T]^T$$

$N!$ possibilities

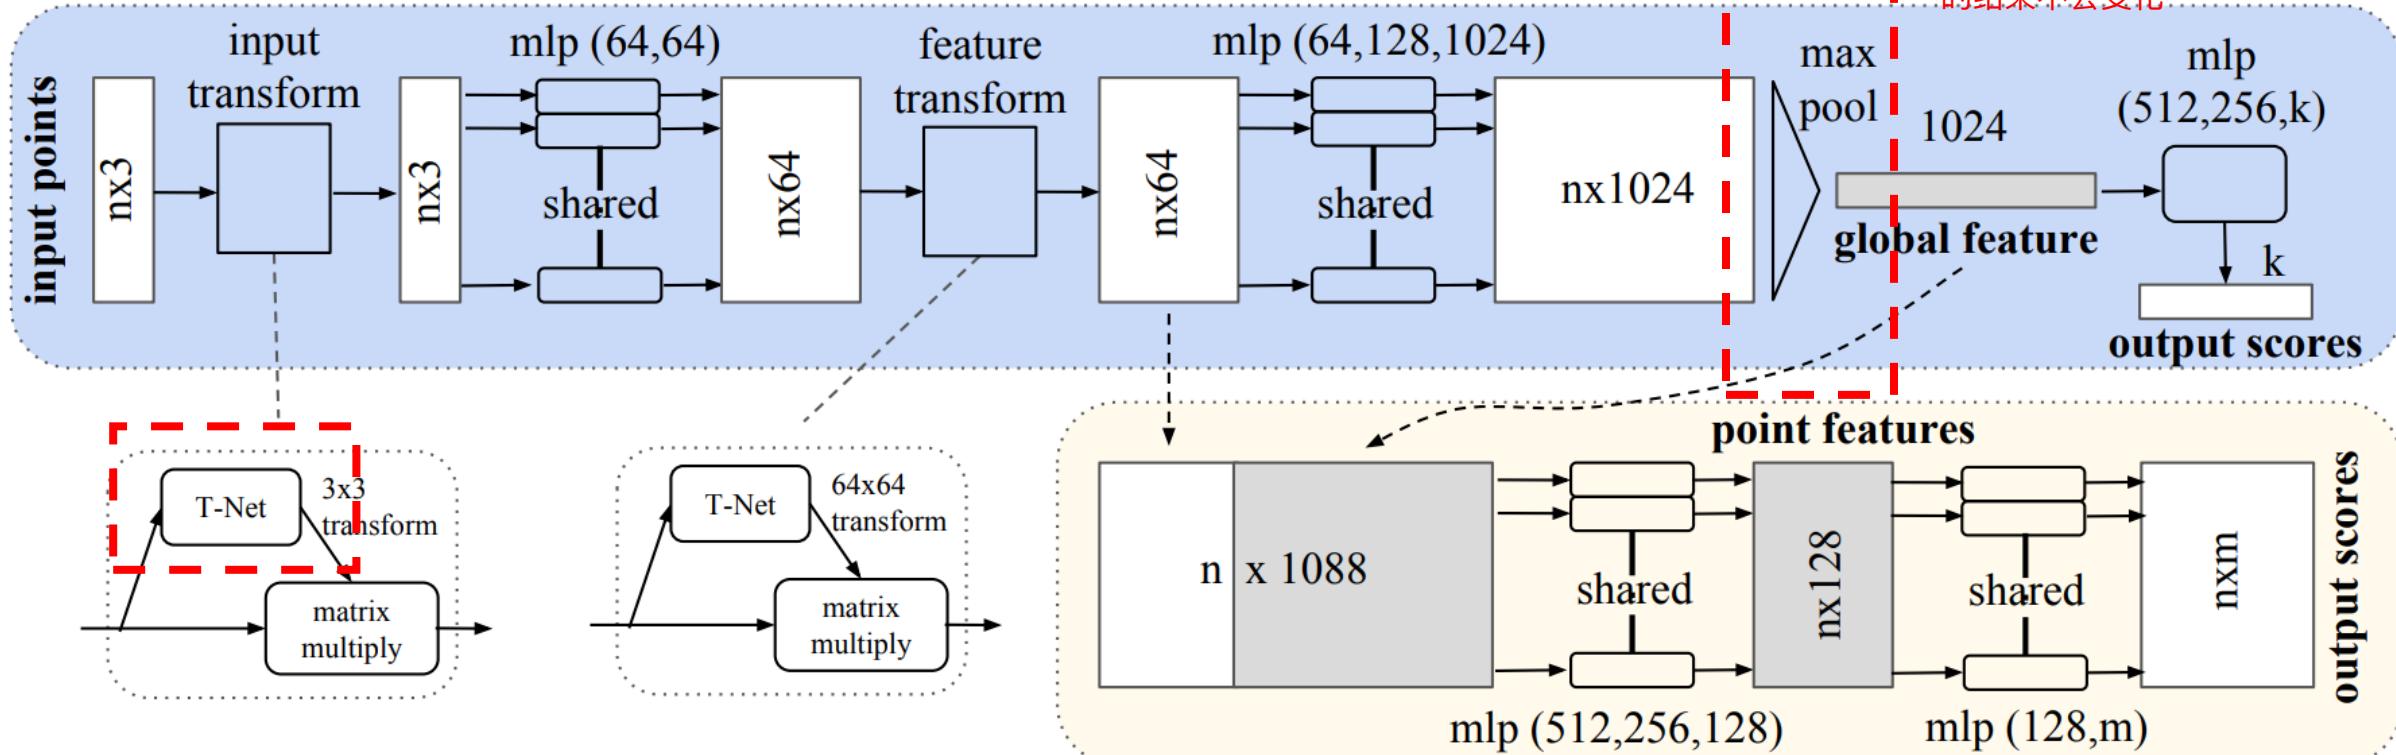




PointNet Architecture

share MLP + Max pool

Classification Network



T-Net is a PointNet itself

想找到一个旋转去得到标准角度

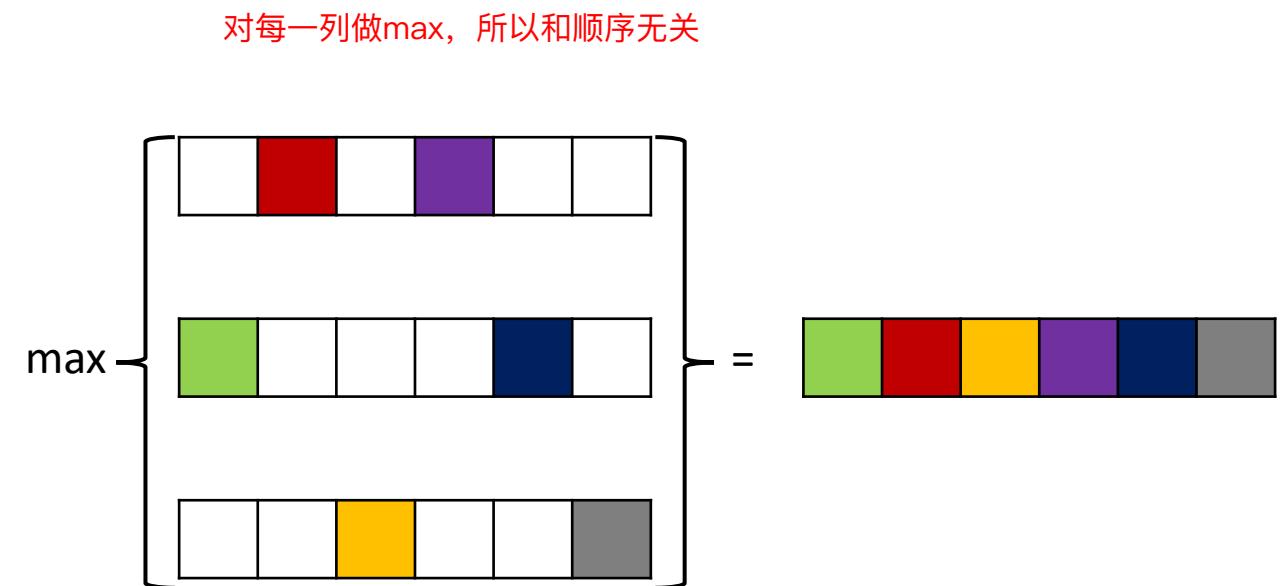
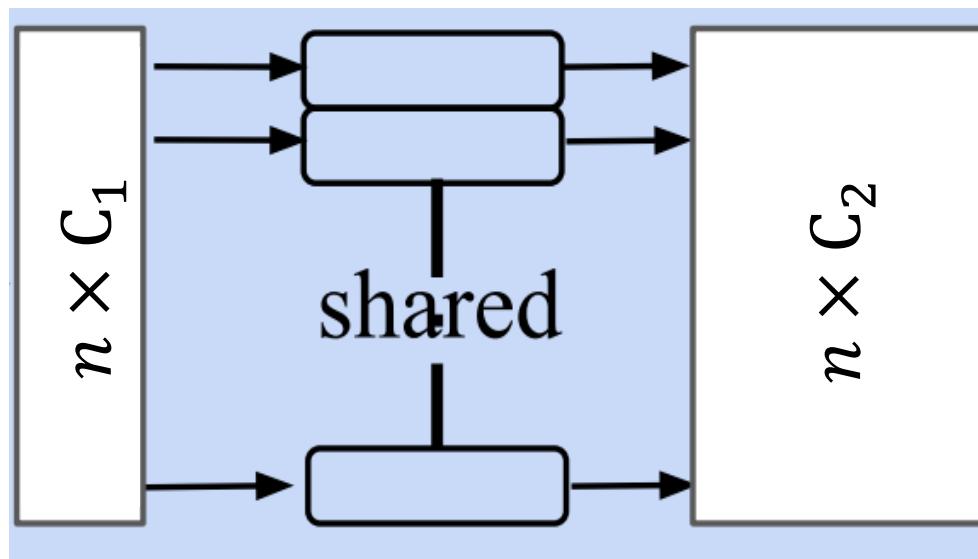
但实际上没有什么用

Segmentation Network

PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation. CR Qi et.al.



- Process each point (feature) **independently** $n \times C_1 \rightarrow n \times C_2$
- Use **Max/Average** to pool the features $n \times C \rightarrow 1 \times C$





◆ Proof – PointNet is able to simulate any function on the point cloud

◆ Some concepts:

- Input points $S = \{x_1, \dots, x_n\}, x_i \in \mathbb{R}^m, x_i \in [0, 1]$.
 - Denote the space of S as χ , i.e., $S \in \chi$ 2或3 都其他 归一化
- A continuous function $f: \chi \rightarrow \mathbb{R}$ 标量
 - This is the “any function” we want to simulate
- *MAX* function: takes n vectors, give element-wise maximum



Given continuous $f: \chi \rightarrow \mathbb{R}$

$\forall \epsilon > 0, \exists h: \mathbb{R}^m \rightarrow \mathbb{R}^{m'}, \text{and } \gamma: \mathbb{R}^n \rightarrow \mathbb{R}$

s.t. $\forall S \in \chi, e.g. S = \{x_1, \dots, x_n\}, x_i \in \mathbb{R}^m$

$$\left| f(S) - \gamma \left(MAX(h(x_1), \dots, h(x_n)) \right) \right| < \epsilon$$



MLP for global feature



Shared MLP



$$\left| f(S) - \gamma \left(\text{MAX}(h(x_1), \dots, h(x_n)) \right) \right| < \epsilon$$

• $h(\cdot)$ maps x_i to some deterministic position of a huge vector

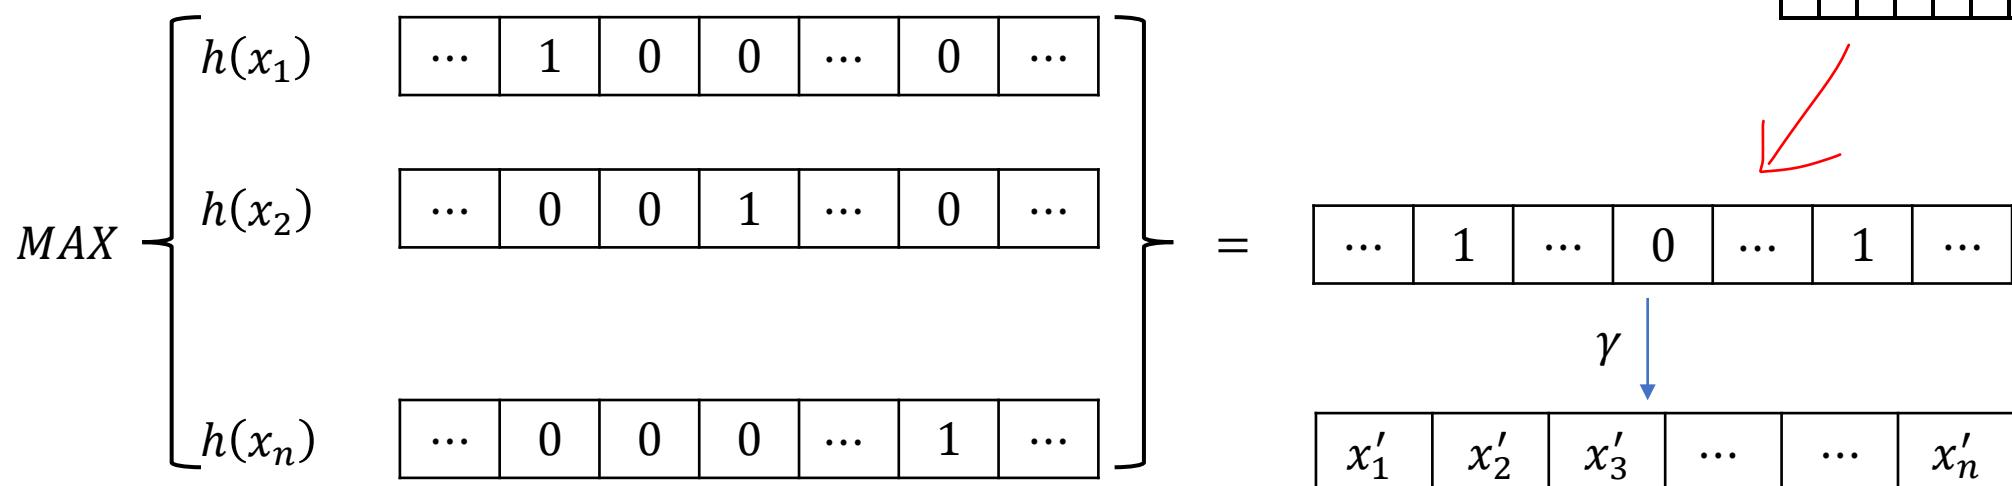
- By **Voxel Grid Downsampling**

• $\text{MAX}(h(x_1), \dots, h(x_n))$ simply builds a voxel grid representation.

- There will be lots of 0 elements because of empty cells in voxel grid.

• $\gamma(\cdot) = \text{reconstruct the points} + f(\cdot)$

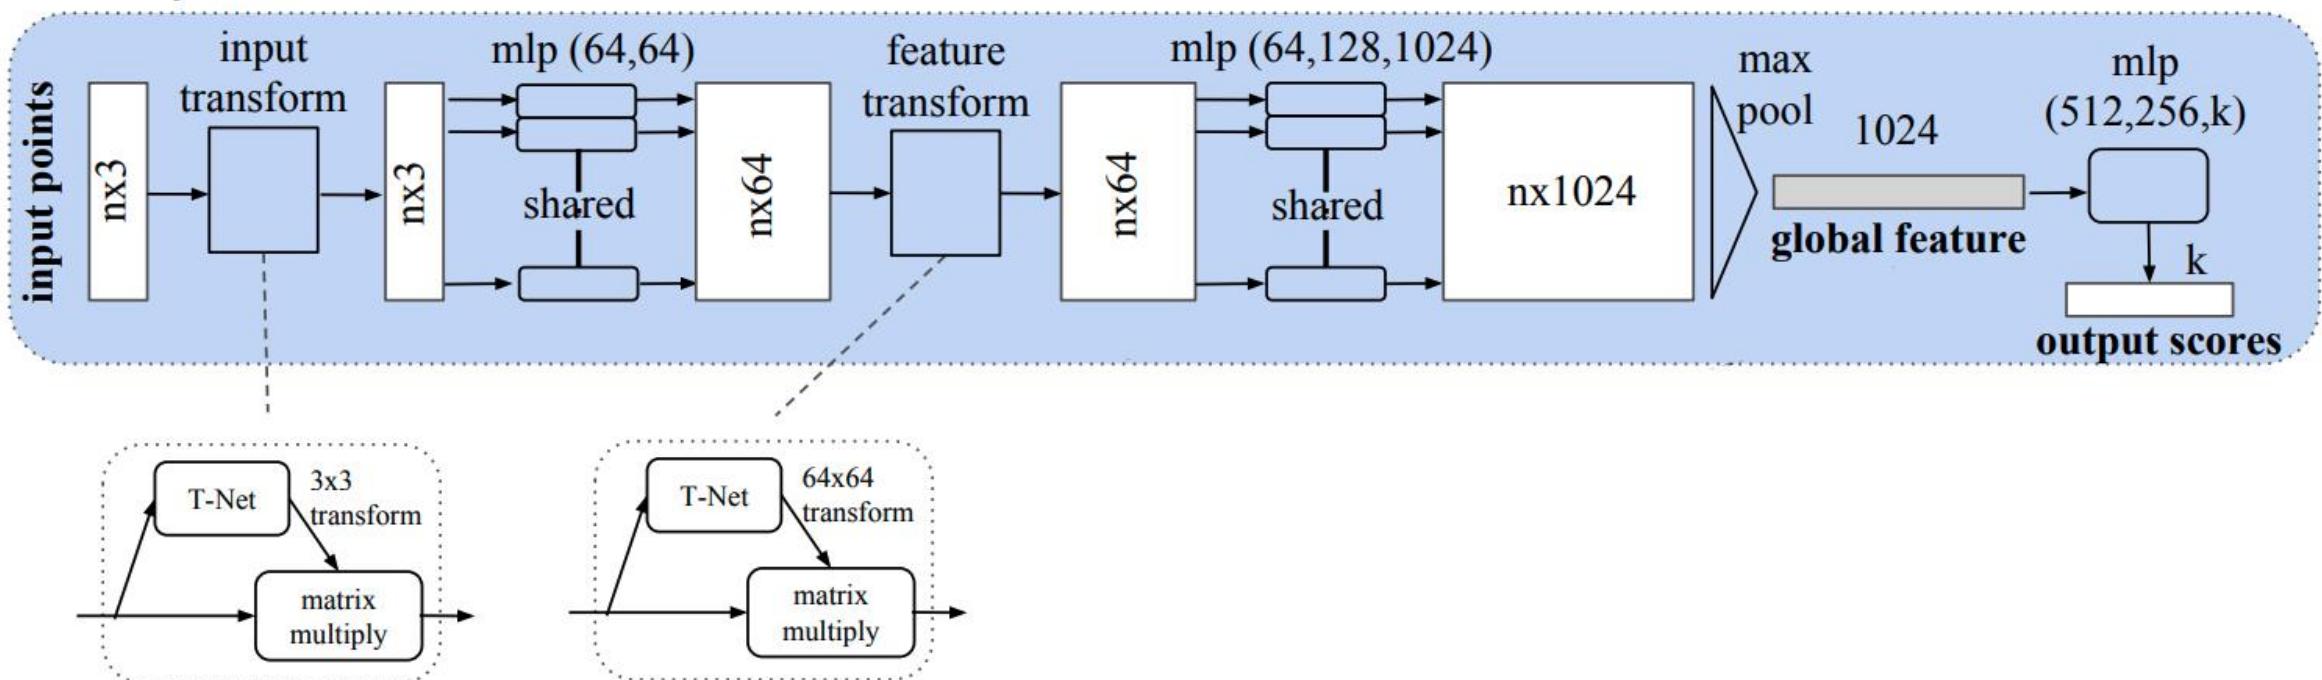
• Done





PointNet – Classification

Classification Network





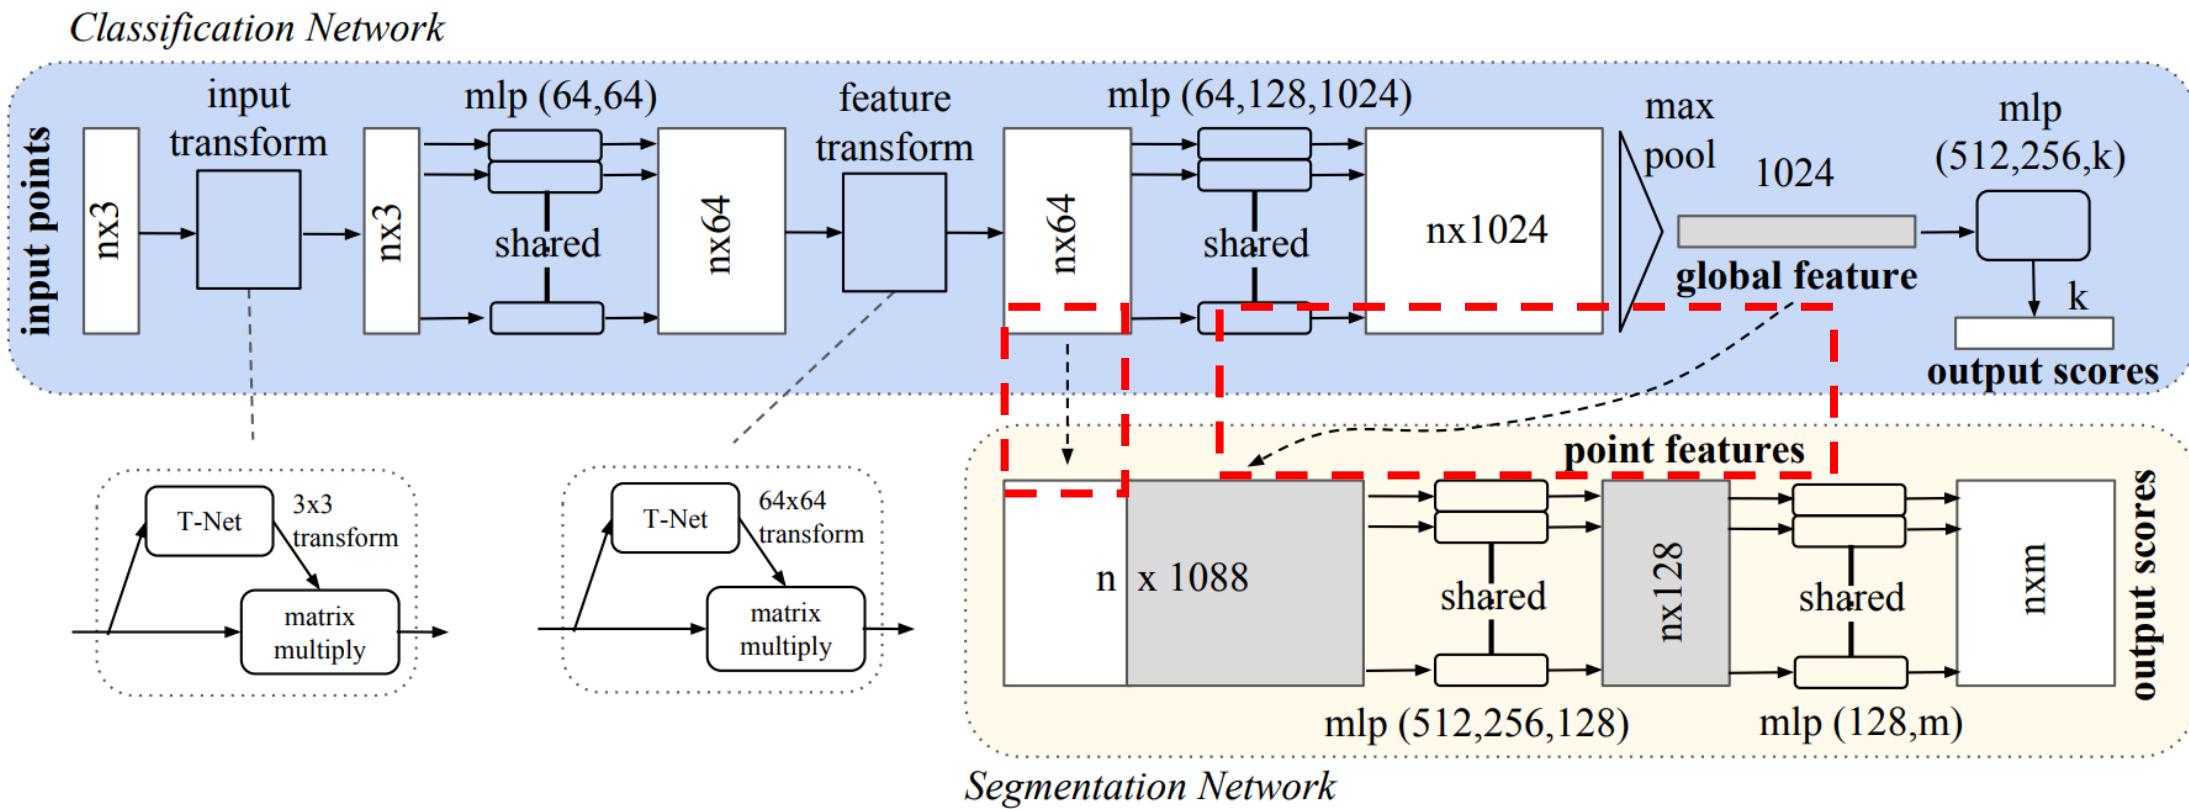
	input	#views	accuracy avg. class	accuracy overall
SPH [11]	mesh	-	68.2	-
3DShapeNets [28]	volume	1	77.3	84.7
VoxNet [17]	volume	12	83.0	85.9
Subvolume [18]	volume	20	86.0	89.2
LFD [28]	image	10	75.5	-
MVCNN [23]	image	80	90.1	-
Ours baseline	point	-	72.6	77.4
Ours PointNet	point	1	86.2	89.2

Table 1. **Classification results on ModelNet40.** Our net achieves state-of-the-art among deep nets on 3D input.



Segmentation is **per-point** classification

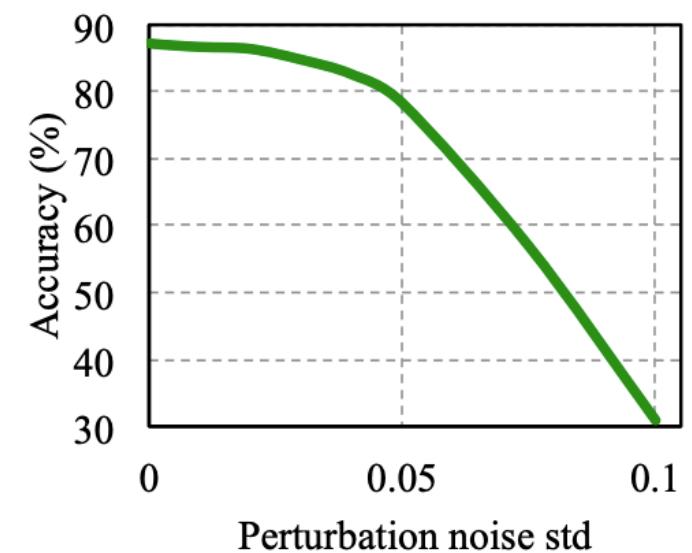
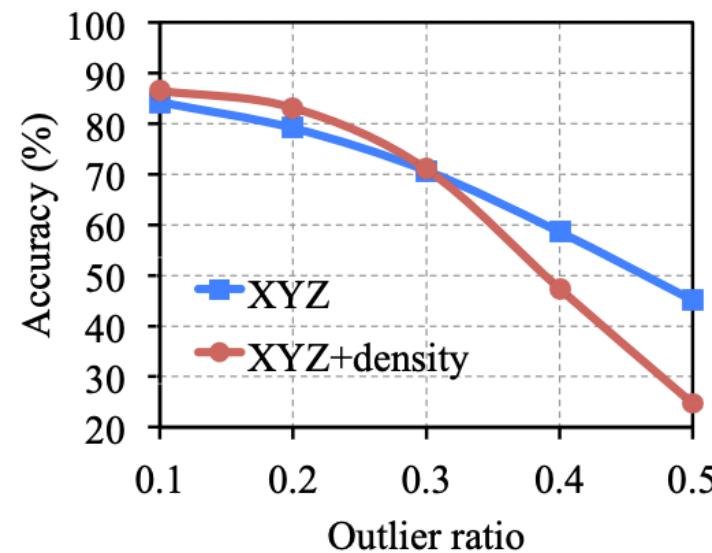
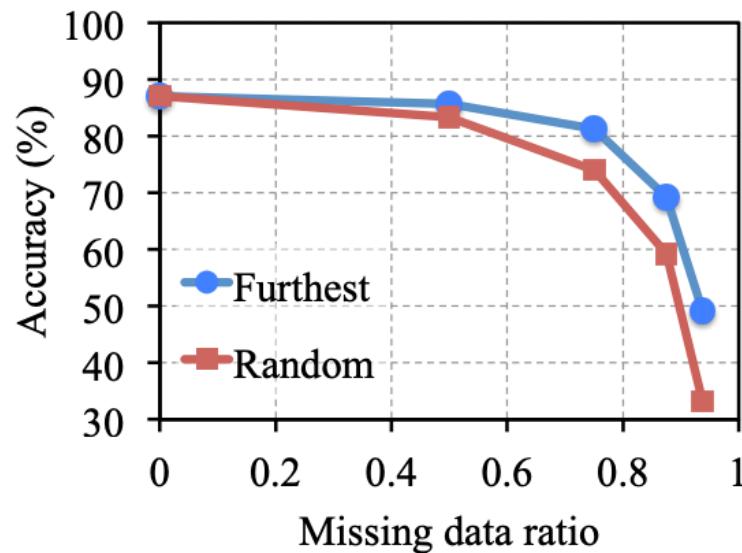
- **MLP on per-point feature**, instead of global feature
- How to get per-point feature?





PointNet – Robustness

robustness 和 max pool 有



Left: Downsample points by random sampling / FPS (farthest point sampling)

Middle: Insert uniformly sampled points.

Right: Add Gaussian noise to input points

$$[x, y, z] \rightarrow [x + \delta_x, y + \delta_y, z + \delta_z]$$

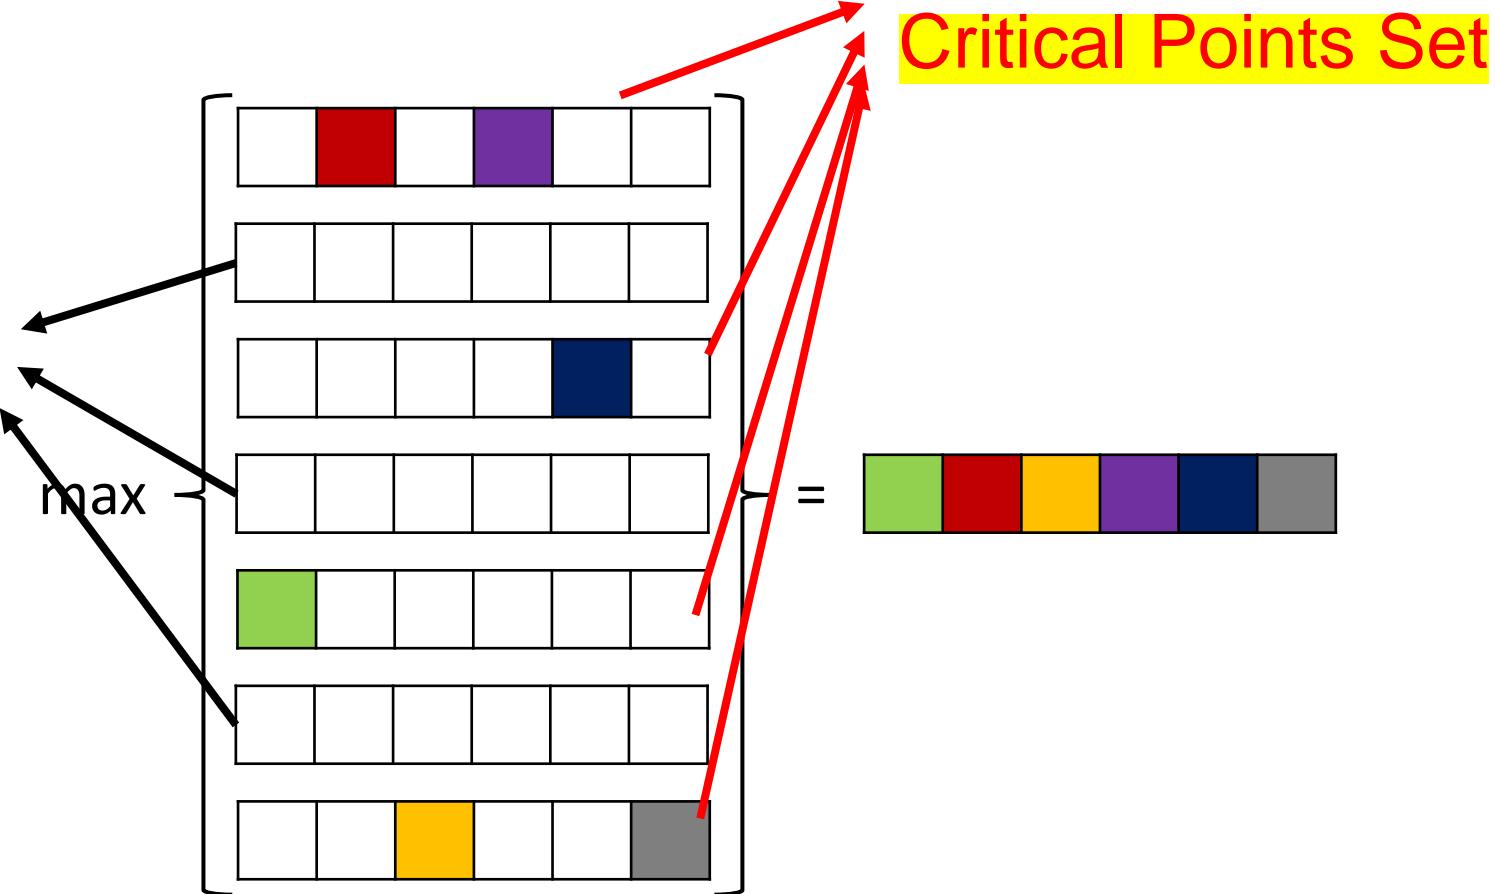


Critical Points Set & Upper Bound Shape



Upper Bound Shape

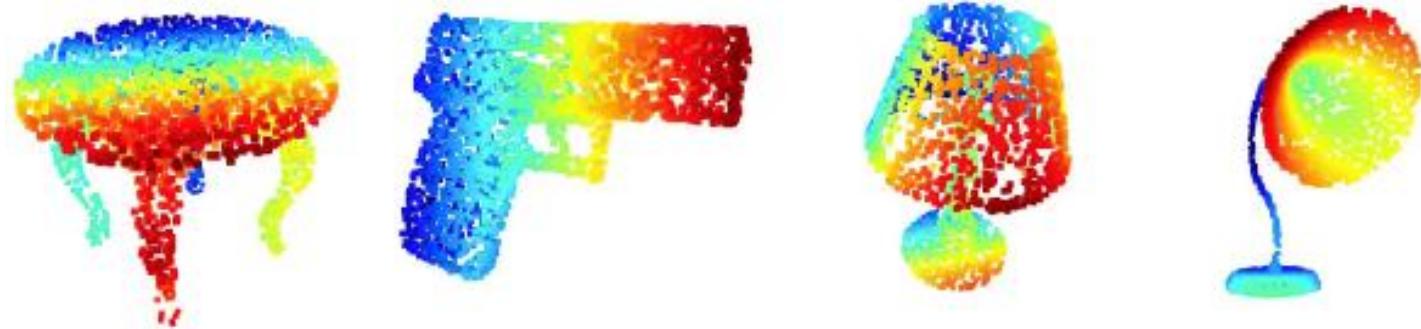
- Add those “useless” points



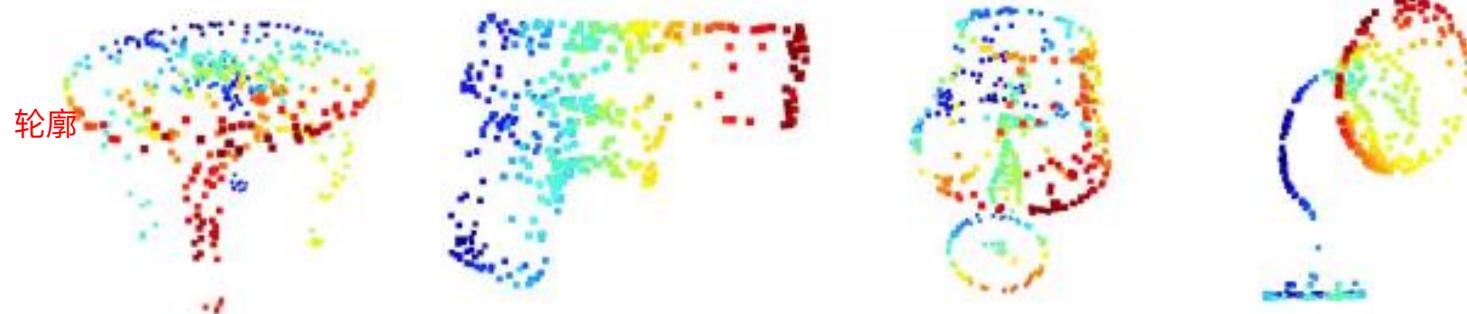


Critical Points Set & Upper Bound Shape

Original Shape

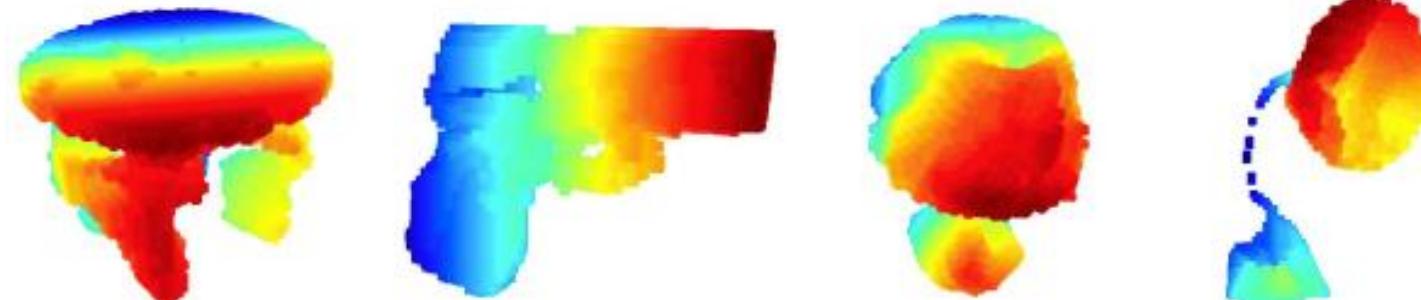


Critical Point Sets



Upper Bound Shape

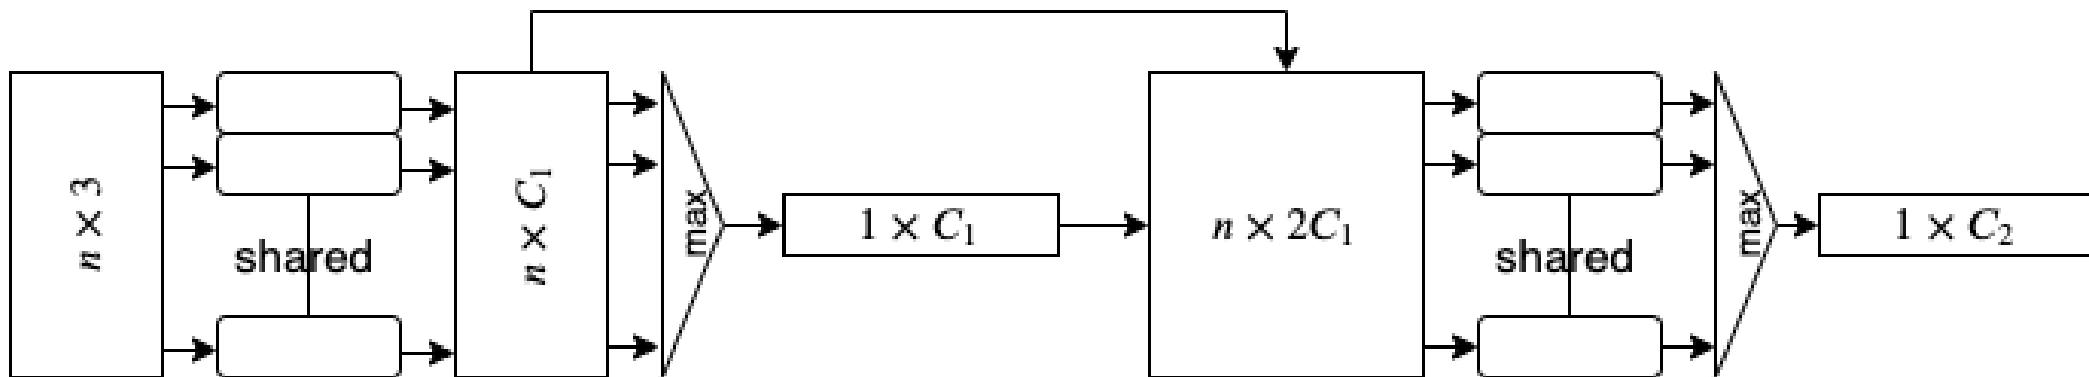
对最终结果不会影响





• Voxel Feature Encoding (VFE)

- Proposed in **VoxelNet**, a **3D object detection** in point cloud, Lecture 6
- Two PN in a row
- Performs better than PN



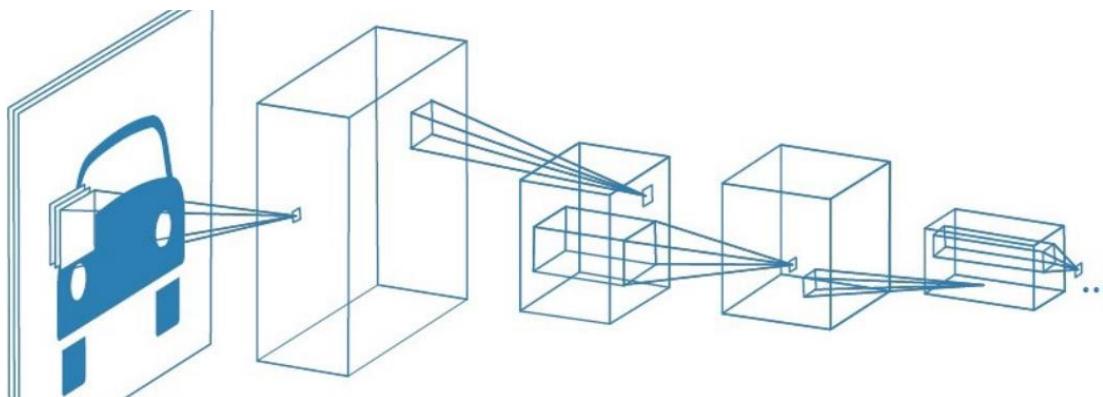


逐层特征提取

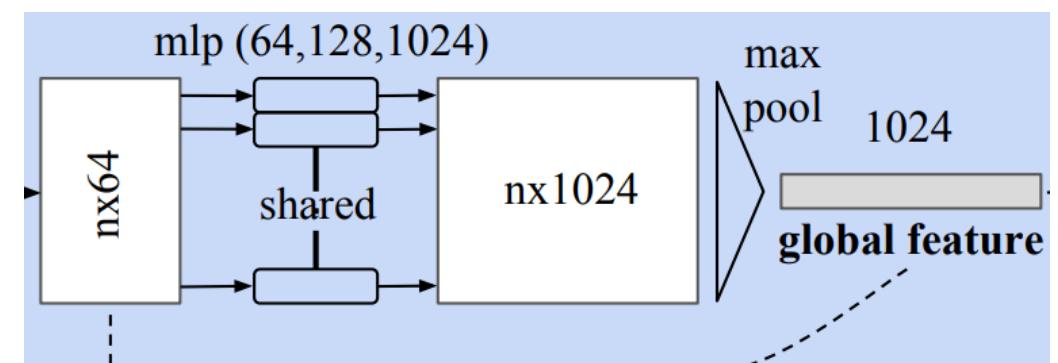
Lack of hierarchical feature aggregation

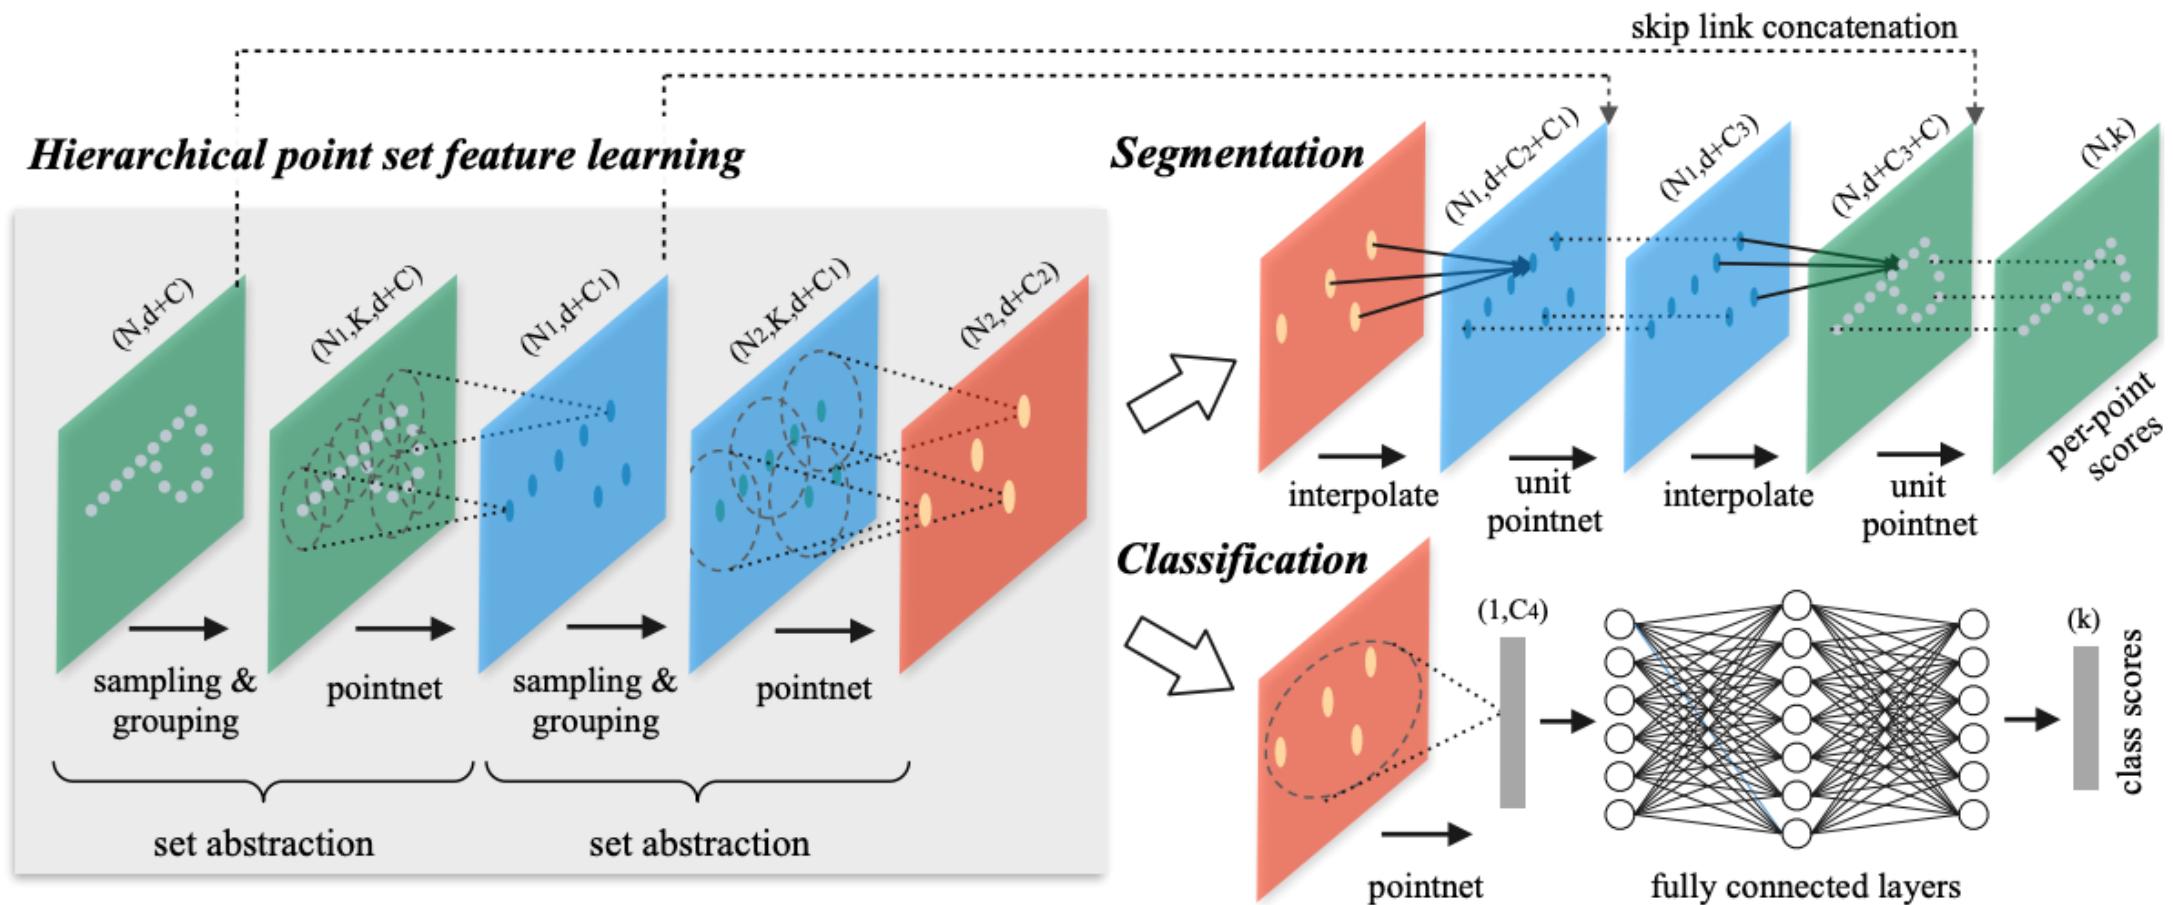
- CNN has multiple, increasing receptive field
- PointNet has one receptive field – all points

CNN



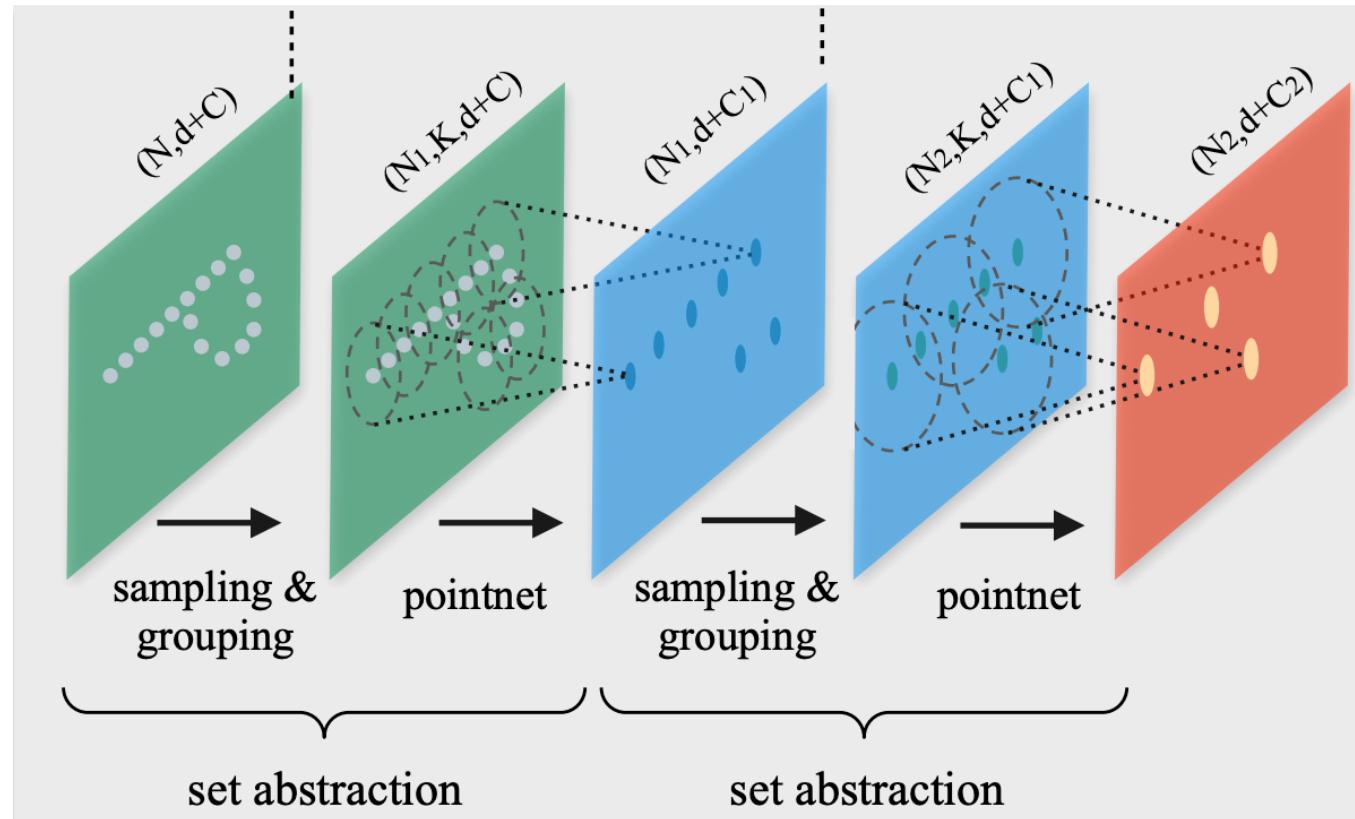
PointNet







PointNet++ - Hierarchical Features



- In each set abstraction:
 - Sampling: FPS
 - Point #: $N_{i-1} \rightarrow N_i$
 - Grouping:
 - Radius Neighbors + random sampling
 - K Nearest Neighbors
- PointNet
 - Point #: N_i
 - Channel #: $C_{i-1} \rightarrow C_i$
 - Concatenate with coordinates so
 $d + C_{i-1} \rightarrow C_i$ 相对坐标
 - Normalize point coordinate in the group
 - Centered with the Node



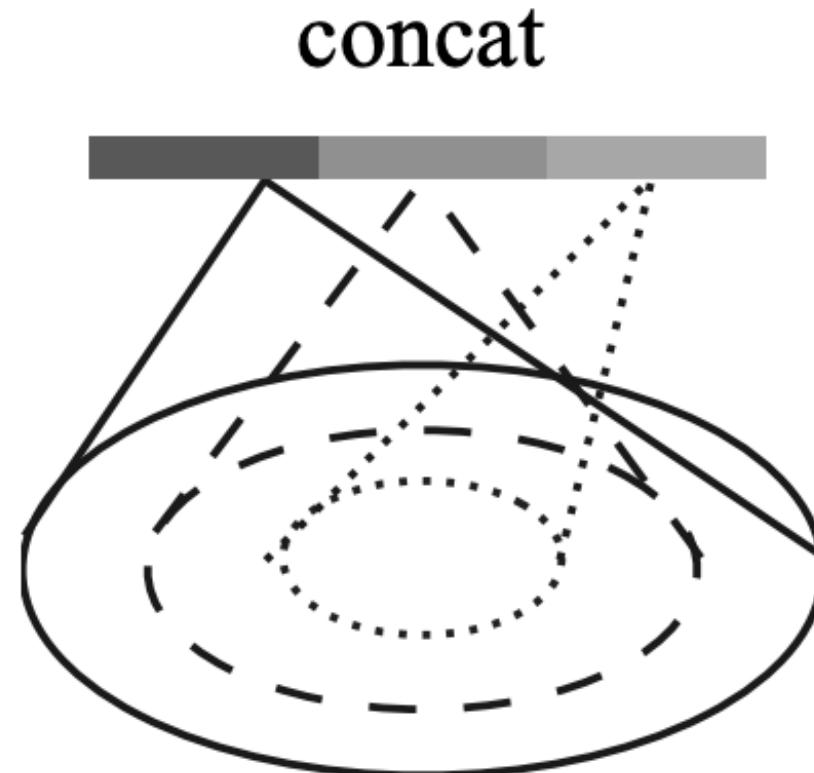
Multi-scale grouping (MSG)

One Sampling

Multiple Grouping & PointNet

- $r = 0.1$ grouping + PN
- $r = 0.2$ grouping + PN
- $r = 0.4$ grouping + PN
- This is compute intensive

Concatenate the multi-scale feature vectors



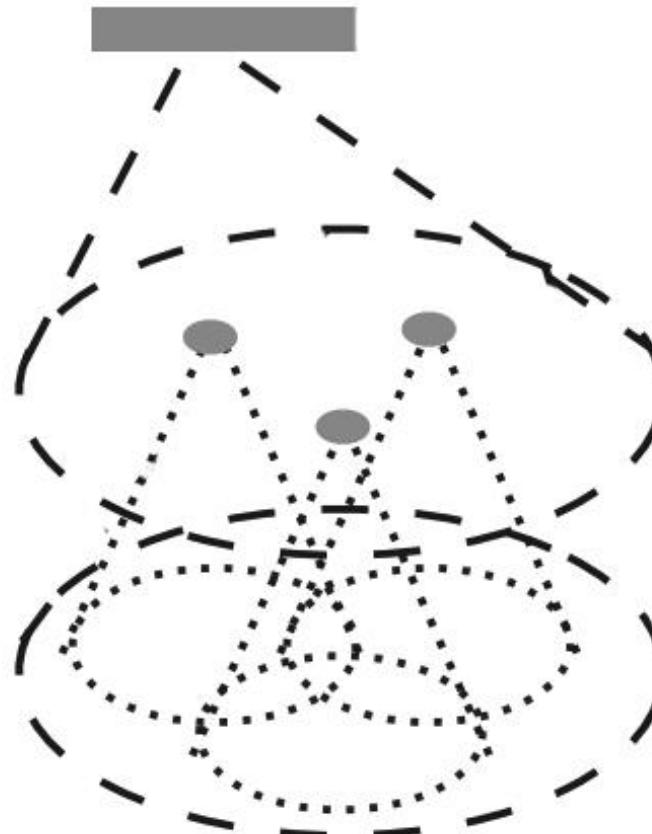


Multi-resolution grouping (MRG)

Get features from

- previous level
- **previous previous** level
- Still increases compute

Simple PN++

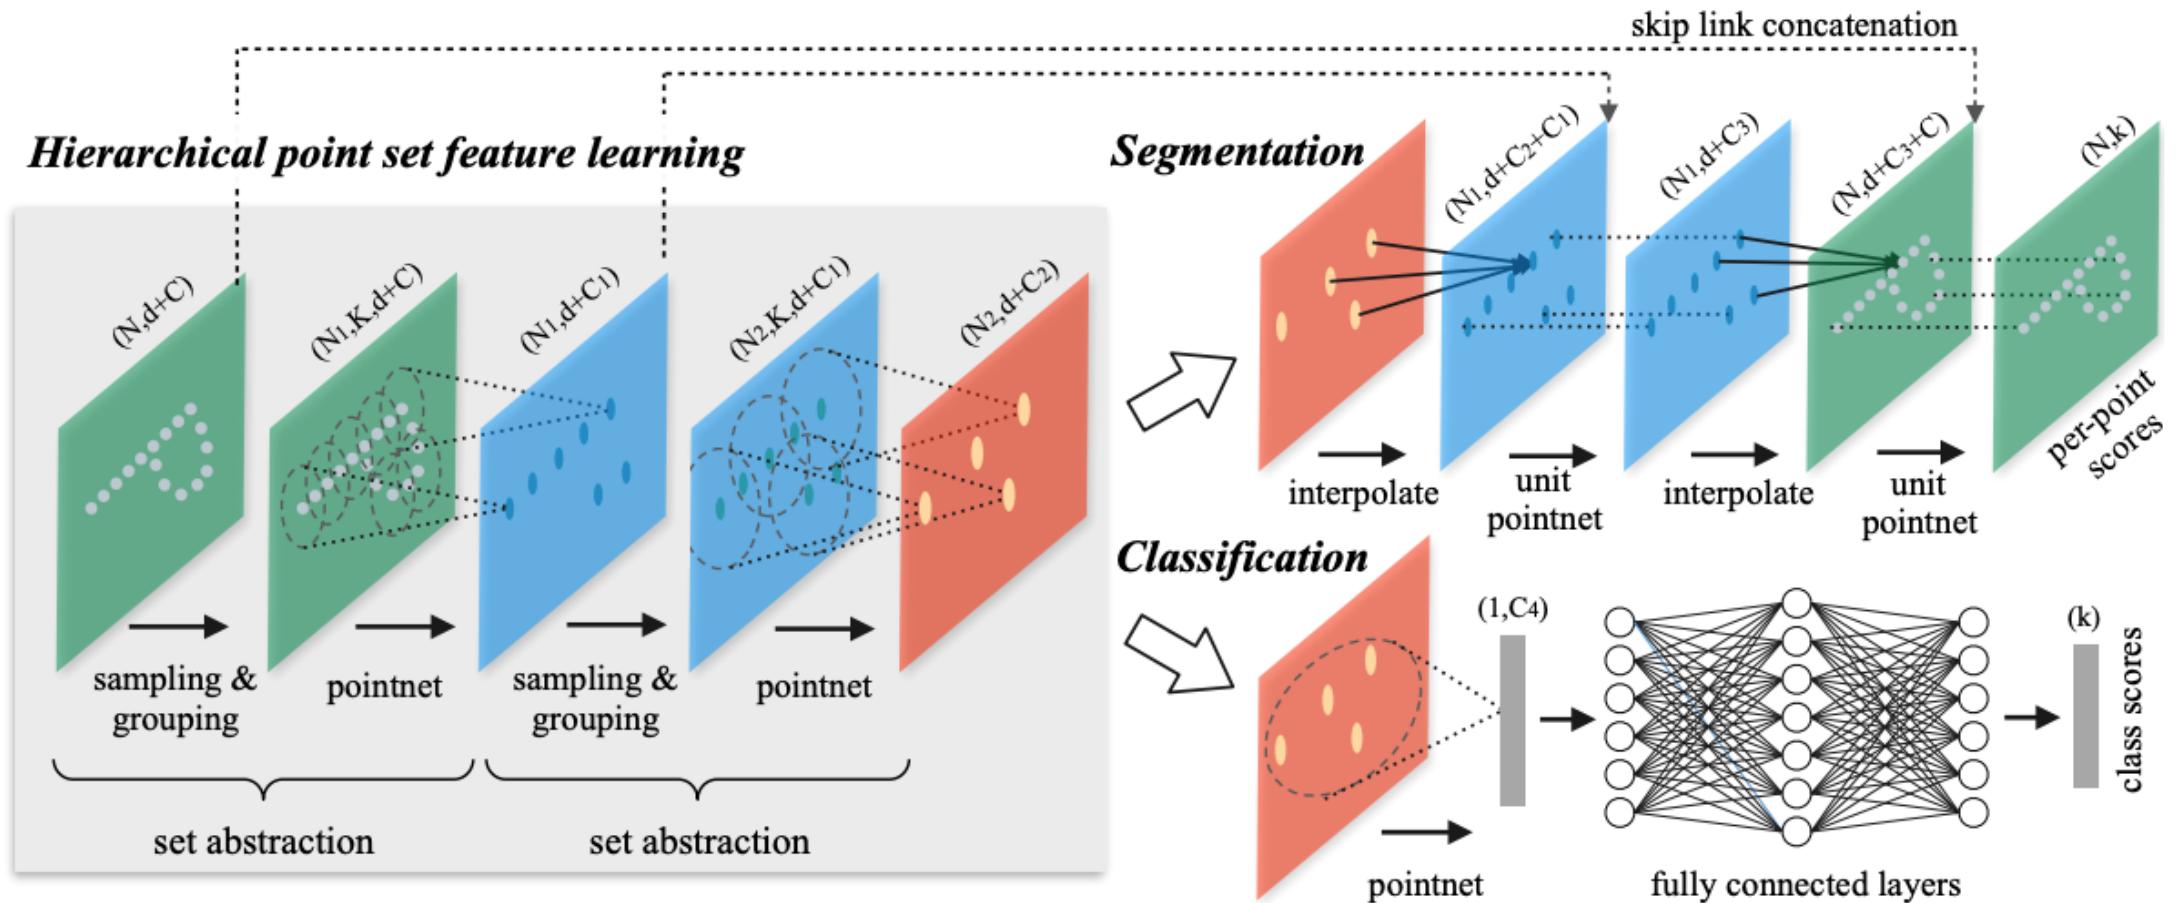


MRG PN++
concat





PointNet++ - Segmentation





Interpolation

- Upsample the features from previous layer

• $x \in \mathbb{R}^3$: point coordinates at the upsampled level, $\# = N_1$

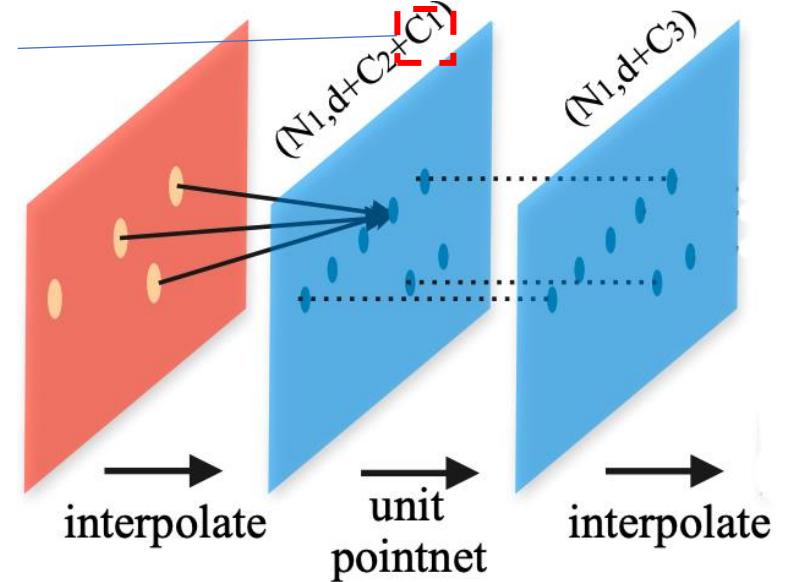
• $f \in \mathbb{R}^{C_2}$: interpolated features

• $x_i \in \mathbb{R}^3$: point coordinates at the previous level (N_2 points)

• $w_i \in \mathbb{R}$: reciprocal of distance $d(x, x_i)$

• $f_i \in R^{C_2}$: point features at the previous level

From encoding stage



$$f^{(j)}(x) = \frac{\sum_{i=1}^k w_i(x) f_i^{(j)}}{\sum_{i=1}^k w_i(x)} \quad \text{where} \quad w_i(x) = \frac{1}{d(x, x_i)^p}, \quad j = 1, \dots, C$$

Nearest neighbour, 越近权越大

 PointNet vanilla: without T-Net

- Same as a single layer in PN++

 There isn't T-Net in PN++

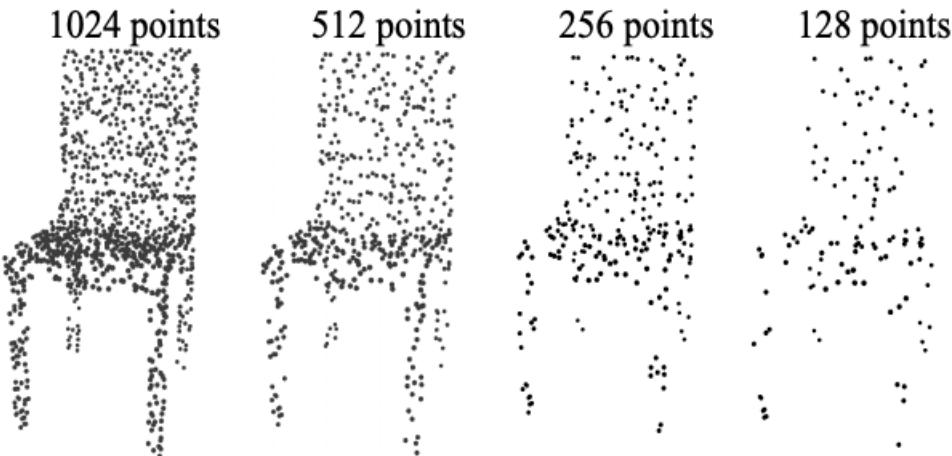
Method	Input	Accuracy (%)
Subvolume [21]	vox	89.2
MVCNN [26]	img	90.1
PointNet (vanilla) [20]	pc	87.2
PointNet [20]	pc	89.2
Ours	pc	90.7
Ours (with normal)	pc	91.9

Table 2: ModelNet40 shape classification.

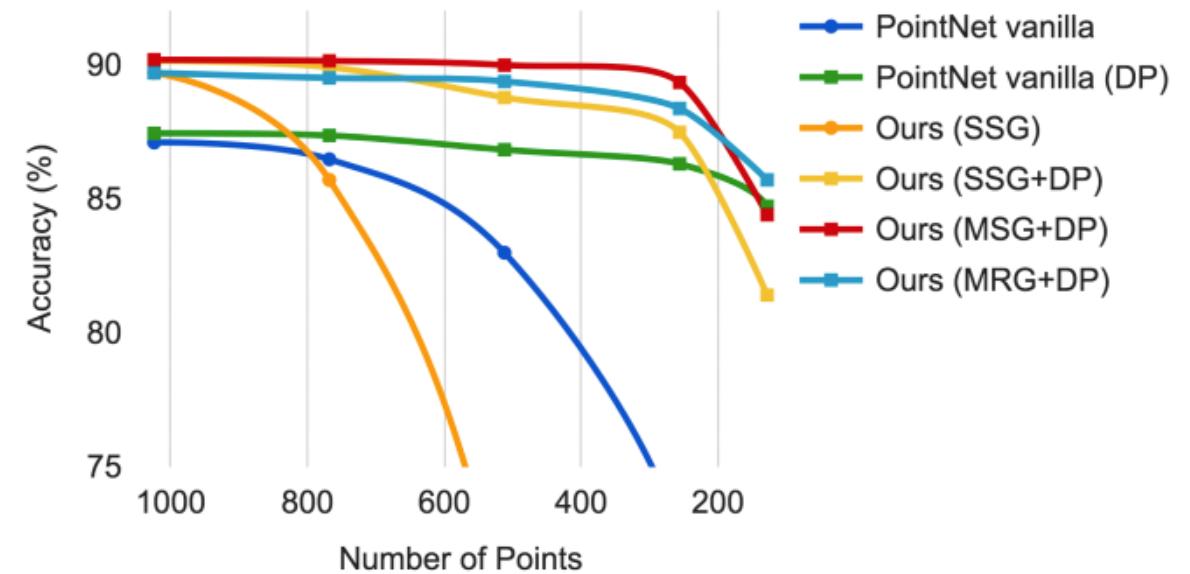


- MSG / MRG collects features from **multiple scale / resolution**
- Performs better with **low intensity**

Random Downsample on ModelNet40



Classification on Downsampled Objects





PointNet

- First method to process point cloud with deep learning
- Lack of hierarchical feature aggregation

PointNet++

- Hierarchical feature aggregation by repeating *sampling-grouping-PointNet*
- Significantly better performance
- Requires slightly more compute



Normalization

- PointNet:
 - Normalized input to zero-mean
- PointNet++:
 - centered-with-node
 - zero-mean

Input Point Dropout

- E.g. Max input point number 5000, randomly dropout to [100, 5000] in each batch

Gaussian Noise

Data augmentation

$$x, y, z + \underline{\Delta x}, \underline{\Delta y}, \underline{\Delta z} \sim N(0, \sigma^2)$$

σ^2 ???

Rotation ???

- Less overfitting
- Worse performance
- Rotation equivariance / invariance is a research topic



- Classification over ModelNet40

- Build your own network with pytorch
 - PointNet example: <https://github.com/fxia22/pointnet.pytorch>
 - ModelNet40 Dataset given by PointNet++:
https://shapenet.cs.stanford.edu/media/modelnet40_normal_resampled.zip
 - Follow the training/testing split
 - Remember to add random rotation over z-axis
 - Report testing accuracy
-
- If you are not familiar with pytorch / deep learning
 - Simply run the open-source code to train / test on the ModelNet40 dataset.
 - Report the testing accuracy.
-
- Please write a report with screenshots to show the training/testing loss/accuracy curve.



Object detection pipeline for lidar

- Use KITTI 3D object detection dataset
- Step 1. Remove the ground from the lidar points
 - Any method you want – LSQ, Hough, RANSAC
- Step 2. Clustering over the remaining points
 - Any method you want
- Step 3. Classification over the clusters
- Step 4. Report the detection precision-recall for three categories: vehicle, pedestrian, cyclist
(Next Lecture)



Step 3. Classification over the clusters

- vehicle / pedestrian / cyclist / other

- Step 3.1. Build dataset from KITTI 3D object detection dataset
 - Extract objects in the box for vehicle / pedestrian / cyclist
 - Other objects? Clustering! generic objects 4类, 建立一个数据集
- Step 3.2. Build deep network to classify them.
- Step 3.3. Report classification accuracy

Step 4. Report the detection Precision-Recall for three categories: vehicle, pedestrian, cyclist

- Fit a cuboid over object detected as vehicle / pedestrian / cyclist
- Generate object detection results for KITTI
- Evaluation code provided next Lecture.